

# Steps into Linear Regression

Quang-Vinh Dinh  
PhD in Computer Science

# Objective

Feature		Label	
	area	price	
	6.7	9.1	
	4.6	5.9	
	3.5	4.6	
	5.5	6.7	

House price data

Features			Label
TV	↕ Radio	↕ Newspaper	↕ Sales
230.1	37.8	69.2	22.1
44.5	39.3	45.1	10.4
17.2	45.9	69.3	12
151.5	41.3	58.5	16.5
180.8	10.8	58.4	17.9

Advertising data

if area=6.0, price=?

if TV=55.0, Radio=34.0,  
and Newspaper=62.0,  
price=?

Features														Label
crim	↕ zn	↕ indus	↕ chas	↕ nox	↕ rm	↕ age	↕ dis	↕ rad	↕ tax	↕ ptratio	↕ black	↕ lstat	↕ medv	
0.00632	18	2.31	0	0.538	6.575	65.2	4.09	1	296	15.3	396.9	4.98	24	
0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.9	9.14	21.6	
0.03237	0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94	33.4	
0.06905	0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.9	5.33	36.2	
0.08829	12.5	7.87	0	0.524	6.012	66.6	5.5605	5	311	15.2	395.6	12.43	22.9	

Boston House Price Data

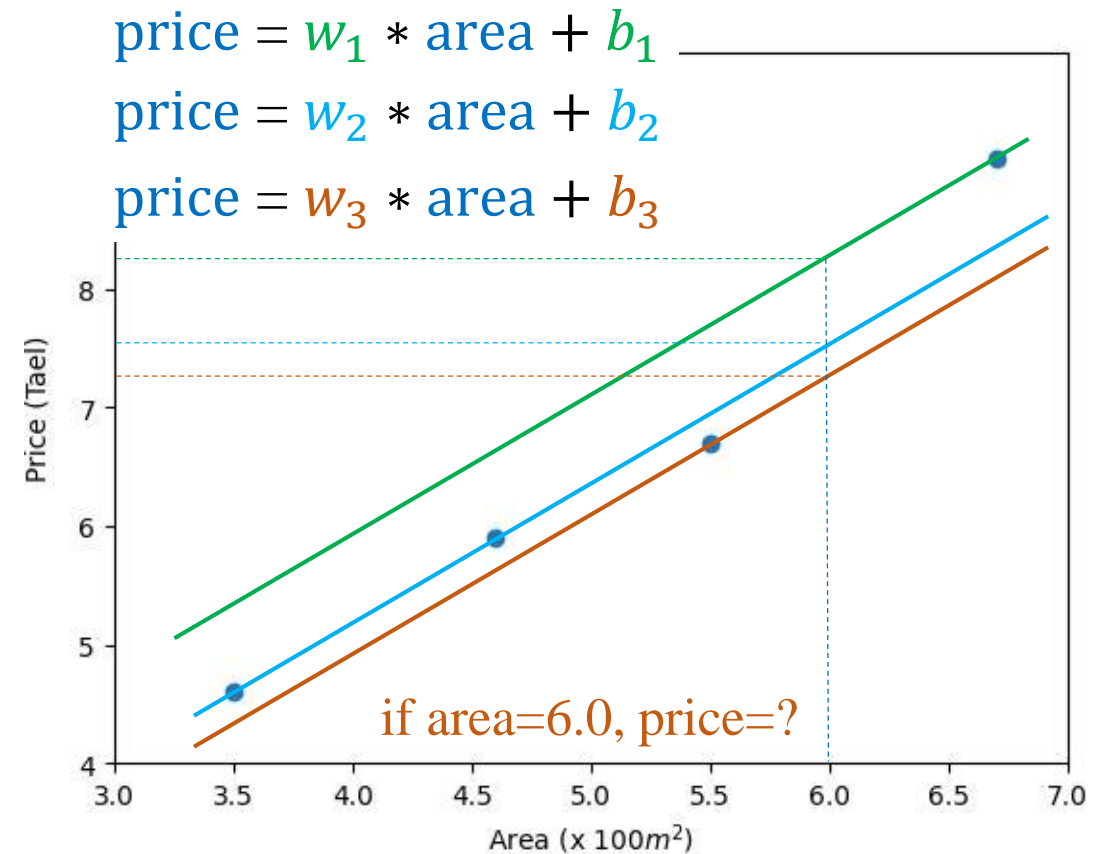
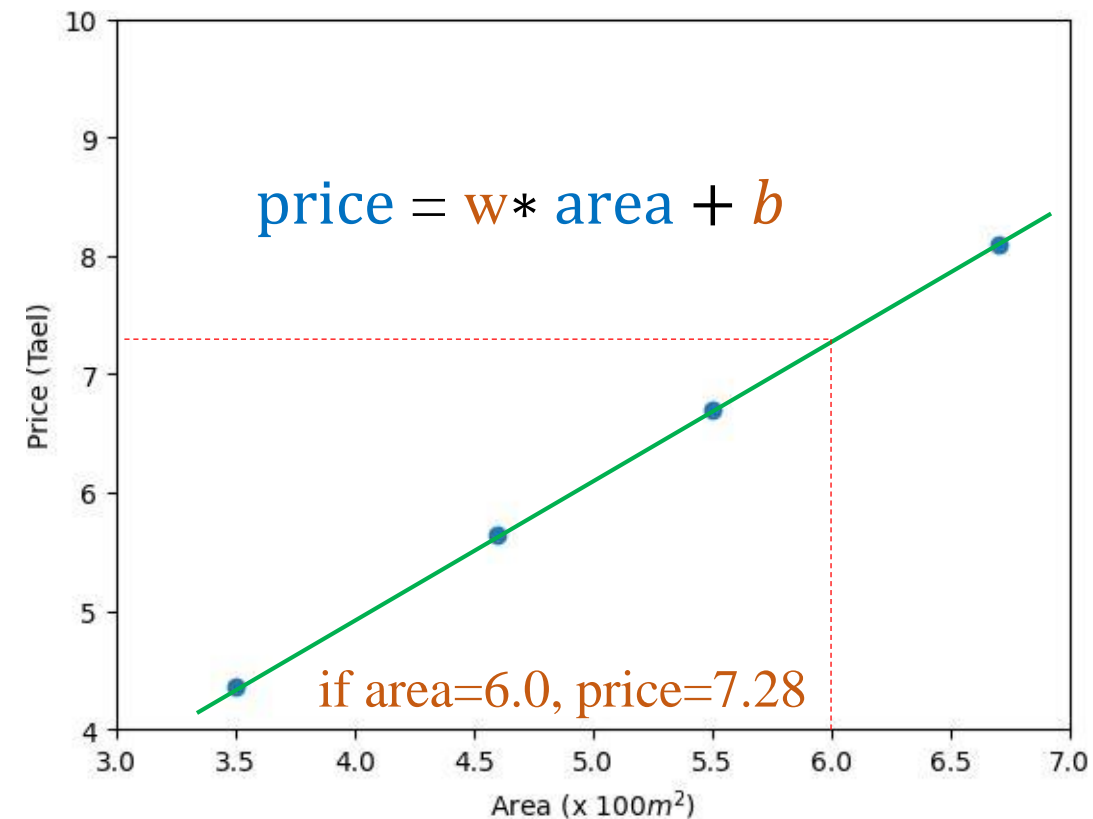
# House Price Prediction

House price data

Feature	Label
area	price
6.7	8.1
4.6	5.6
3.5	4.3
5.5	6.7

Feature	Label
area	price
6.7	9.1
4.6	5.9
3.5	4.6
5.5	6.7

House price data

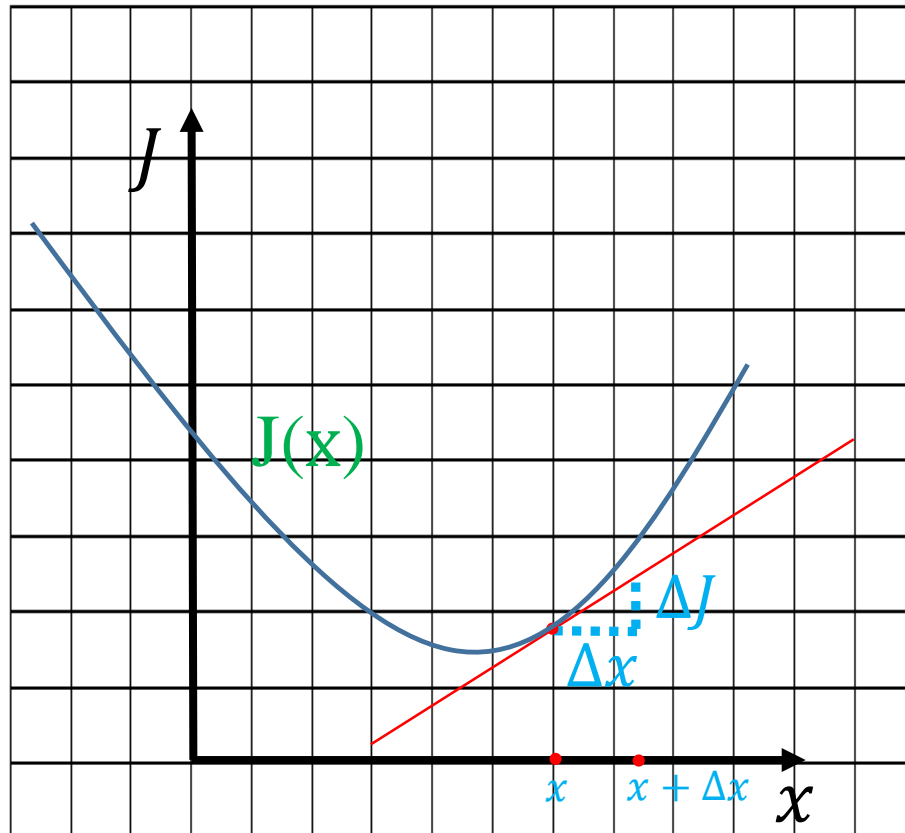


# Outline

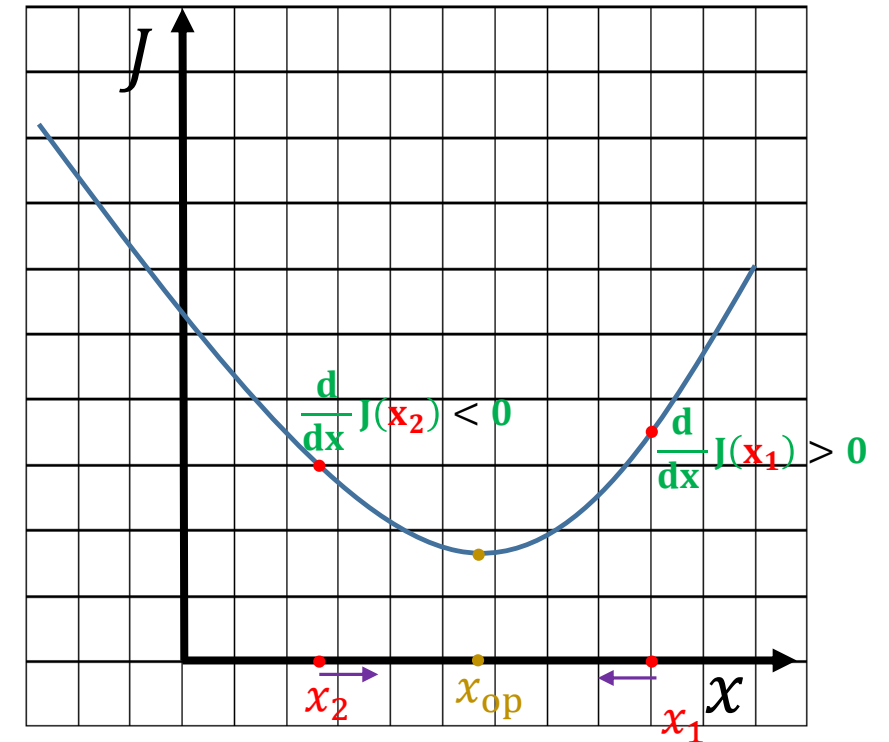
- **Review on Optimization**
- **Partial Gradient and Chain Rules**
- **Finding a Line**
- **Linear Regression**

# Optimization

## ❖ Gradient descent



$$\frac{d}{dx}J(x) = \lim_{\Delta x \rightarrow 0} \frac{J(x + \Delta x) - J(x)}{\Delta x}$$



$$x = x - \eta \frac{d}{dx}J(x)$$

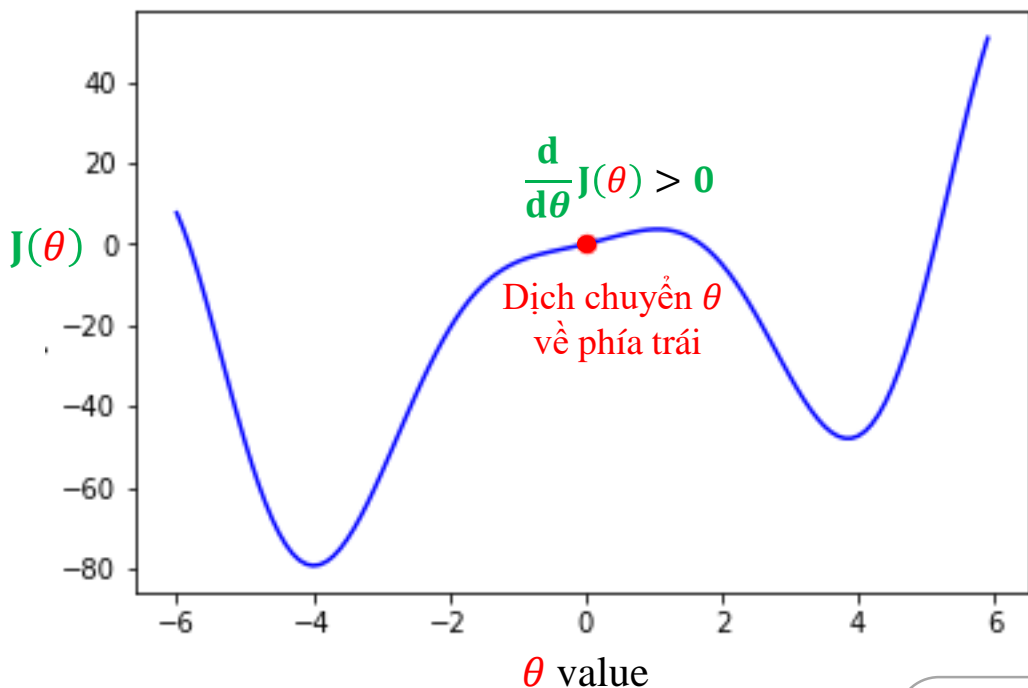
Derivate at x

learning rate

# Optimization

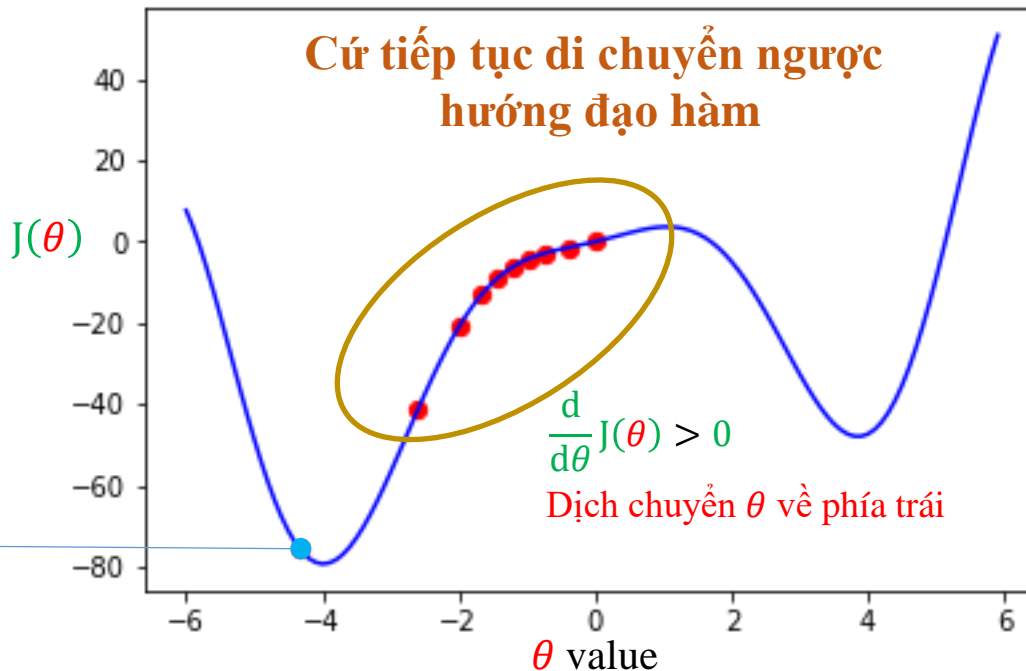
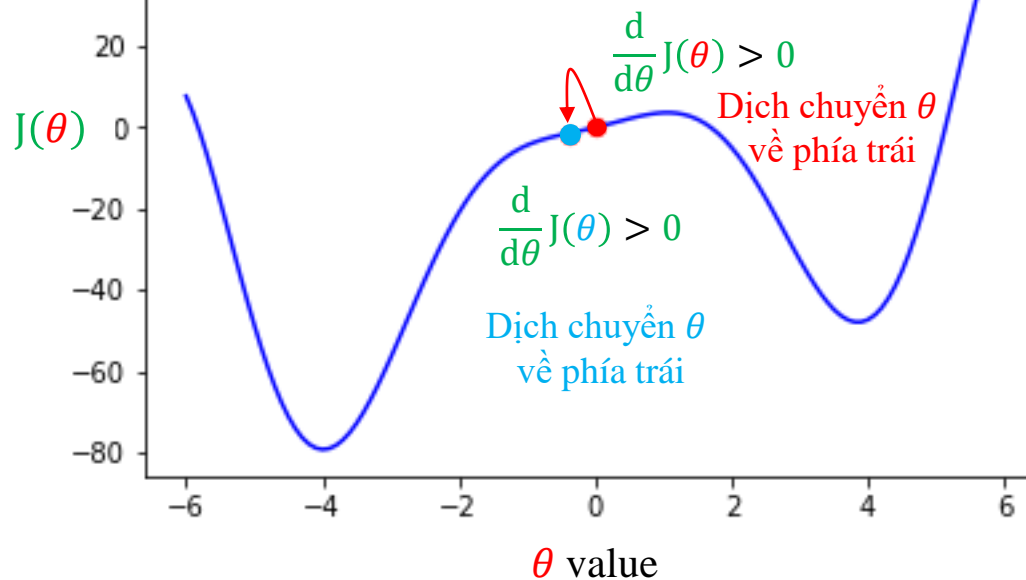
## ❖ Gradient descent

### Khởi tạo giá trị $\theta$



$\frac{d}{d\theta} J(\theta) < 0$   
Dịch chuyển  $\theta$  về phía phải

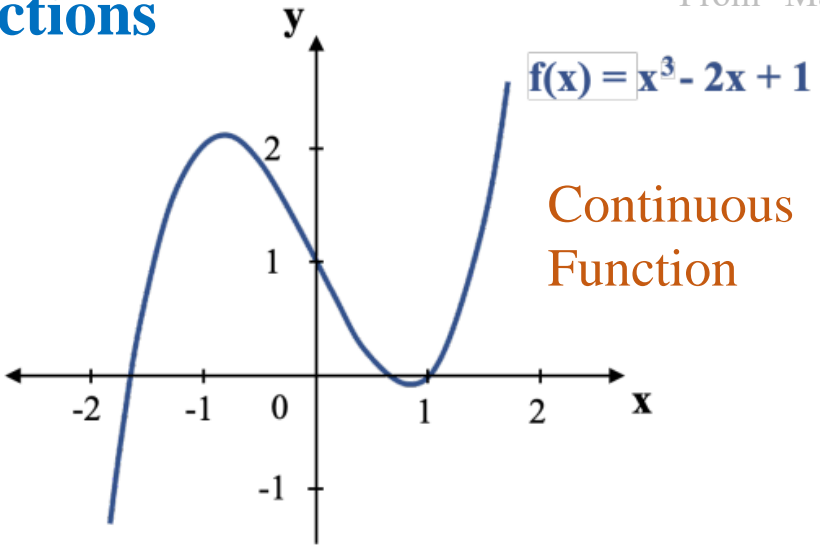
### Di chuyển $\theta$ ngược hướng đạo hàm



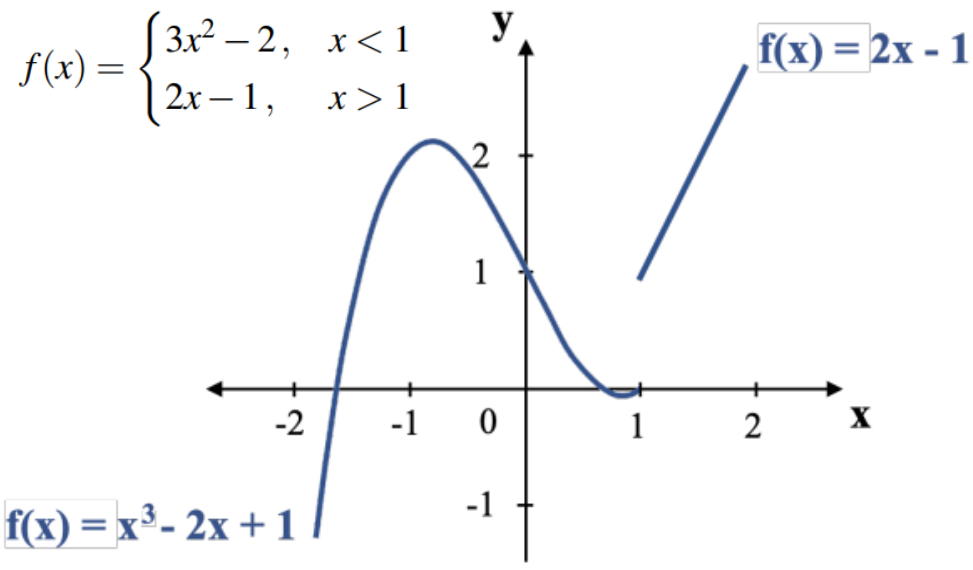
# Optimization Algorithms

## Loss functions

From "Machine Learning Simplified"

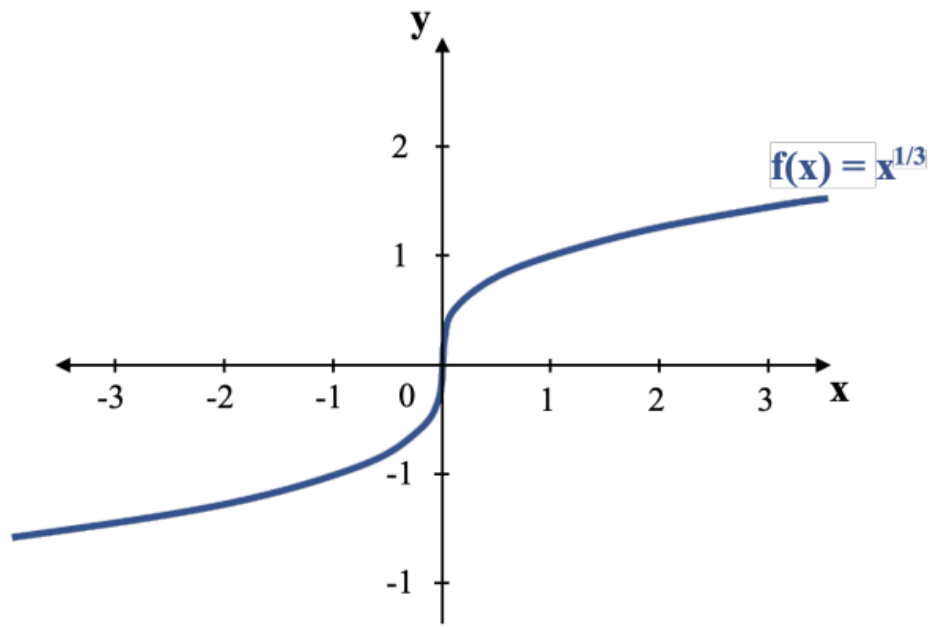


Continuous  
Function

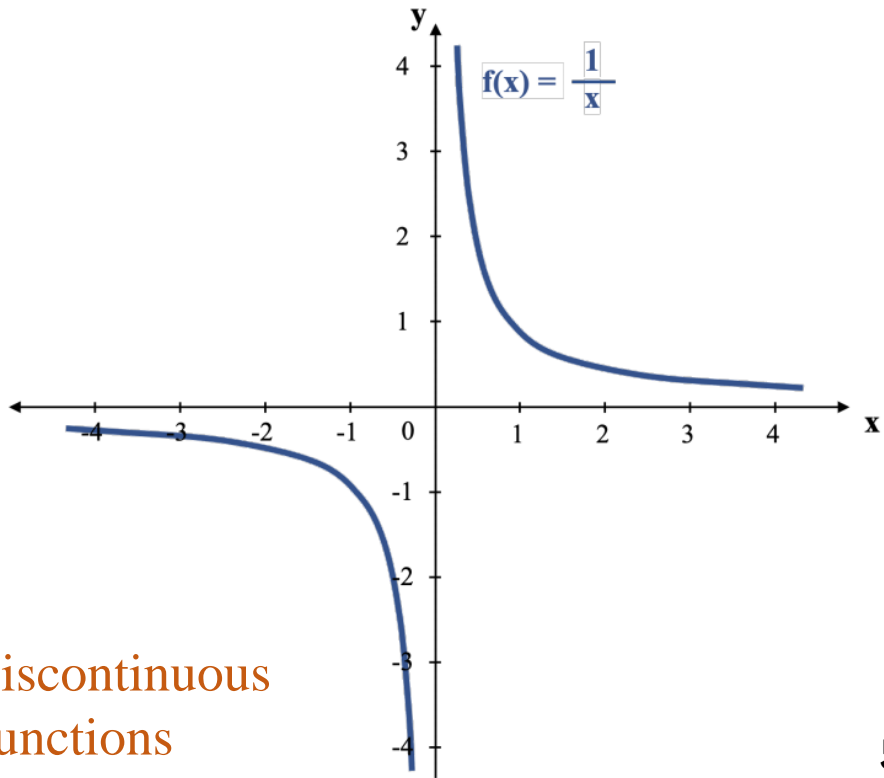


$f(x) = x^3 - 2x + 1$

$f(x) = 2x - 1$



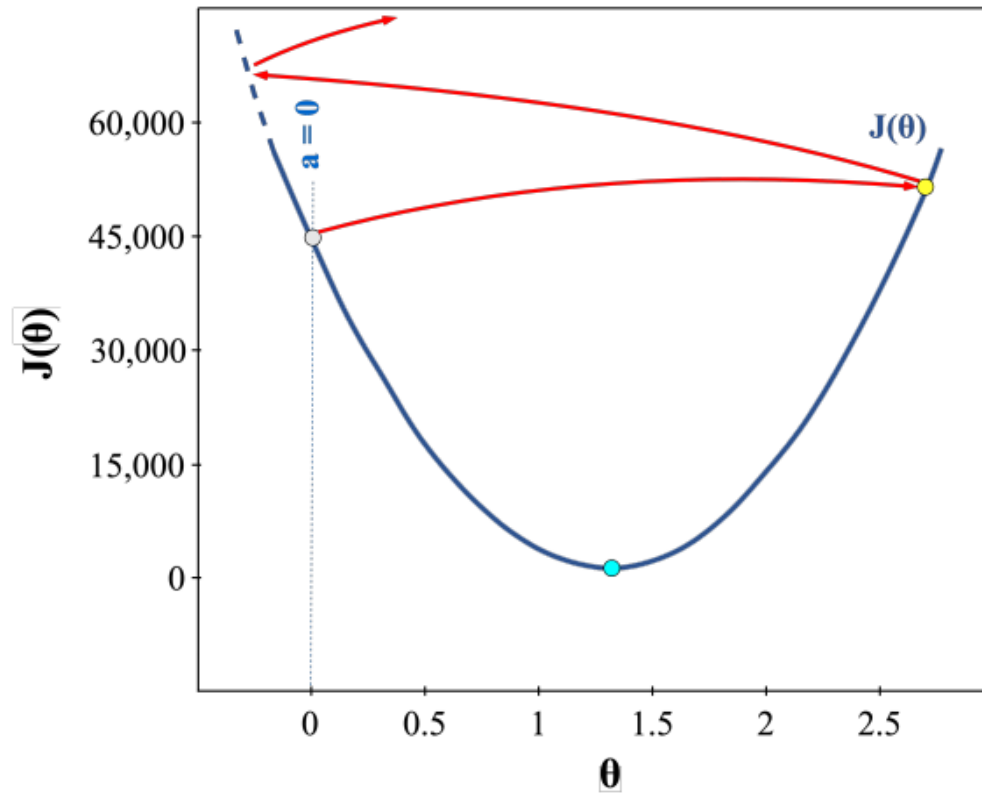
Continuous non-differentiable functions



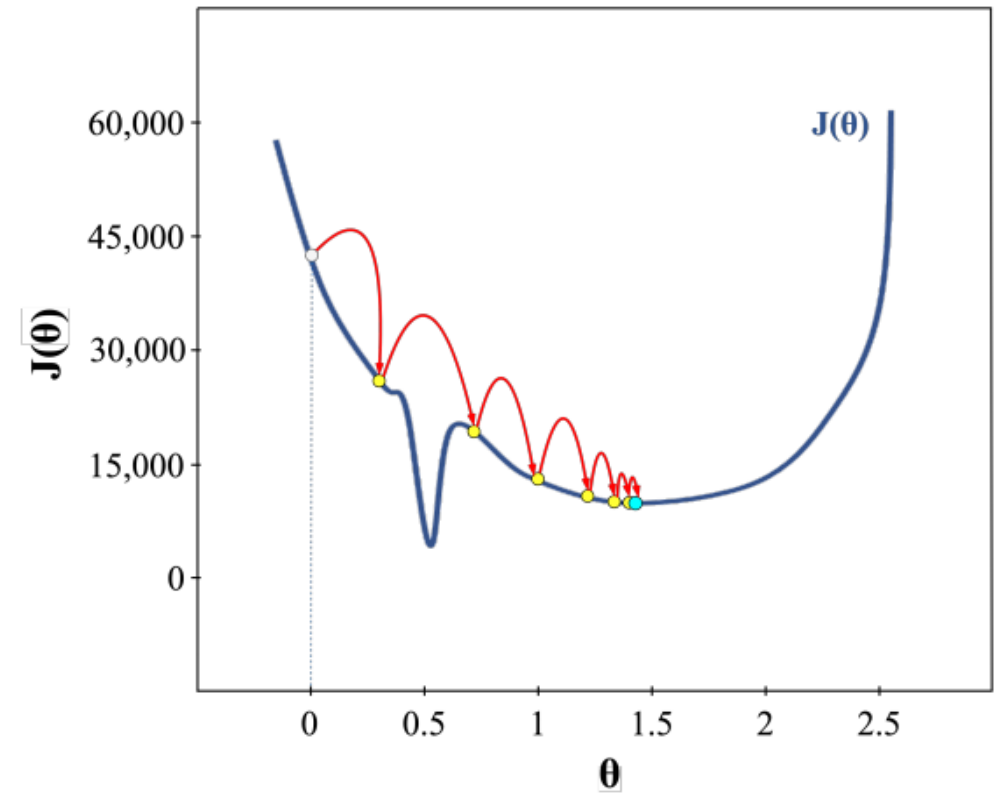
Discontinuous  
Functions

# Optimization Algorithms

## Learning rate



(a) Gradient descent missing global minimum on a convex cost function due to a very large learning rate.



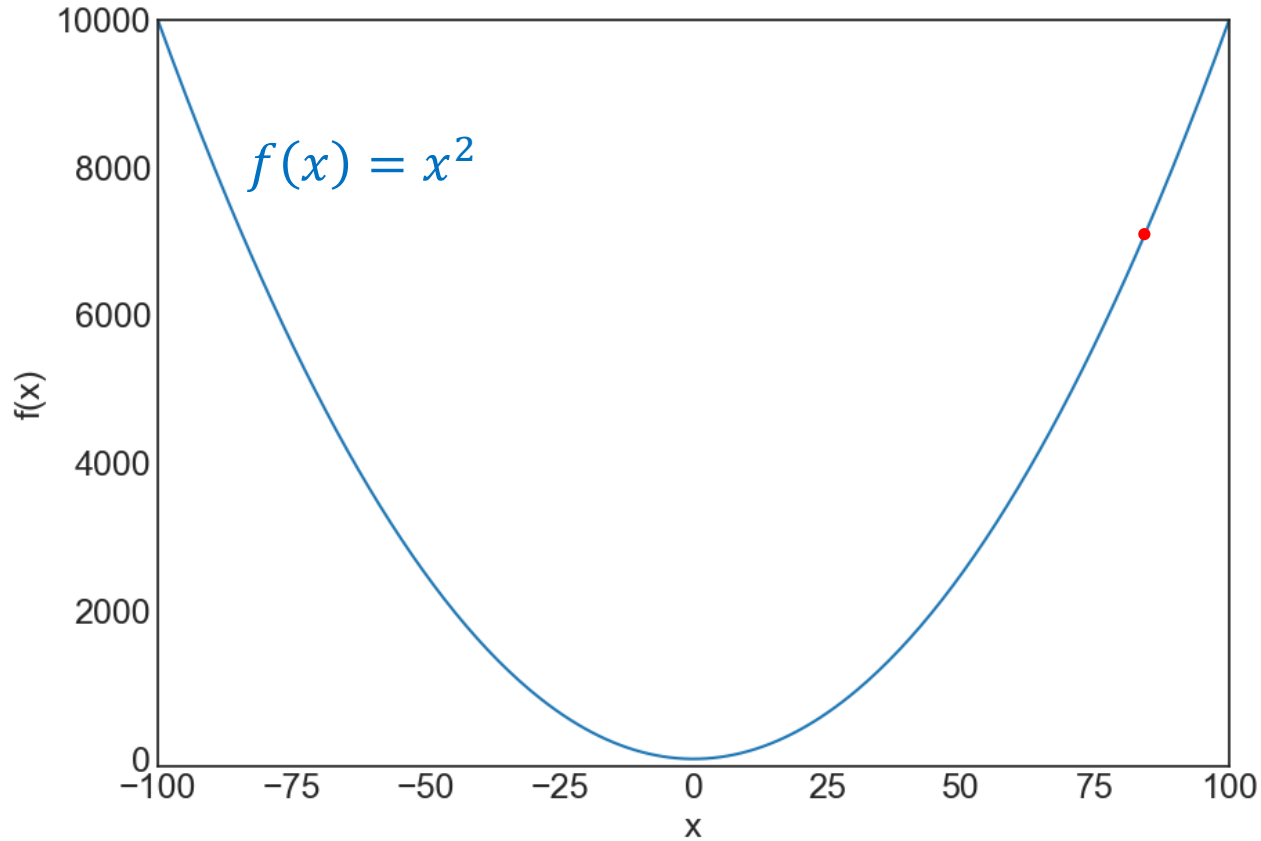
(b) Gradient Descent missing global minimum on a non-convex cost function due to a very large learning rate.





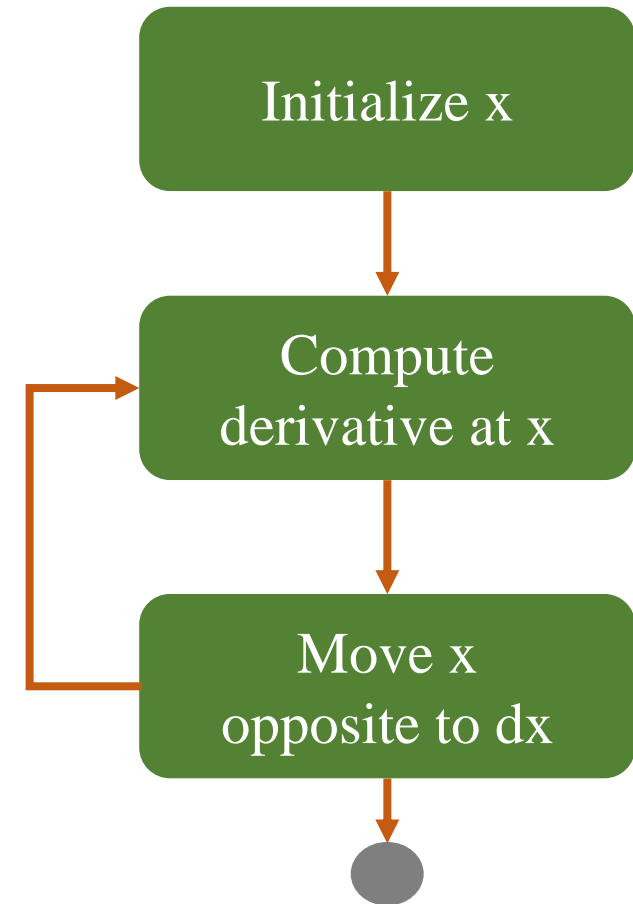
# Gradient-based Optimization

## ❖ Square function



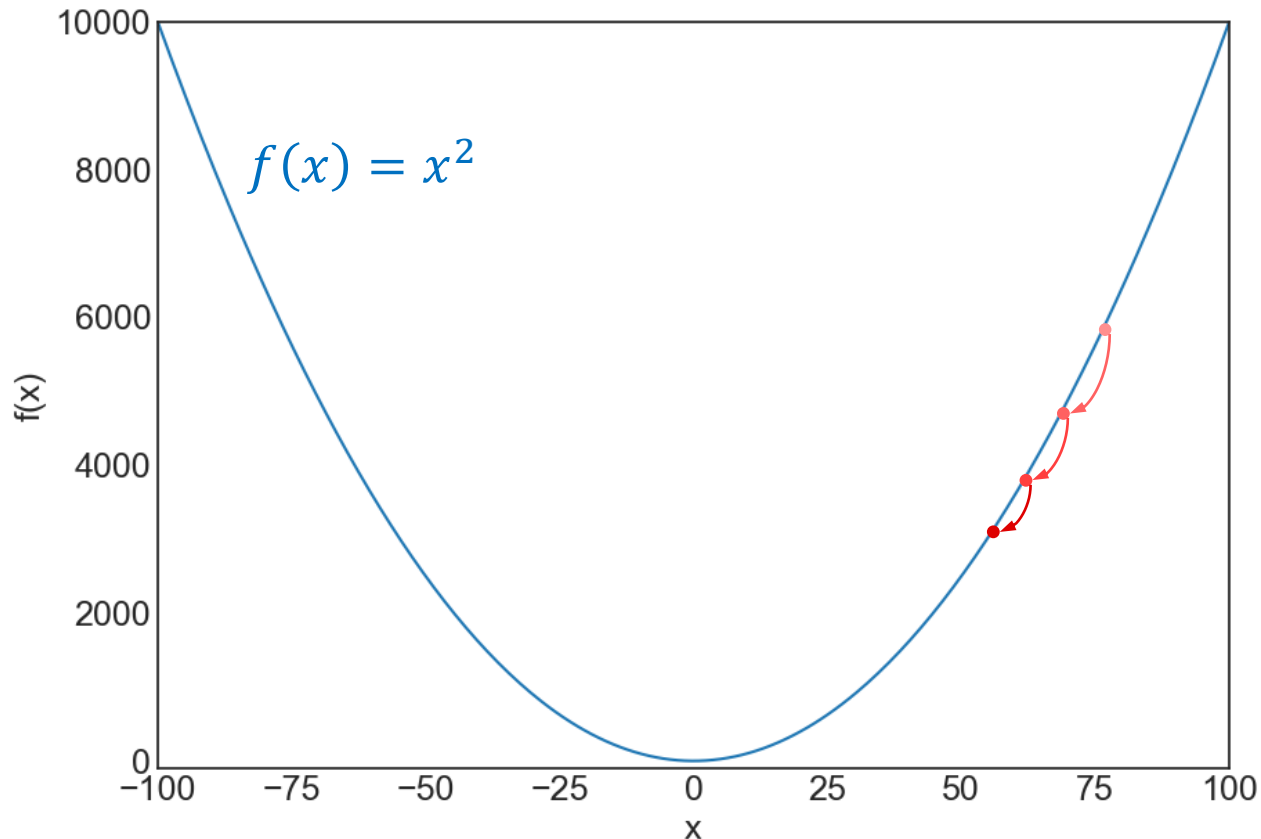
$$\begin{aligned} -100 \leq x \leq 100 \\ x \in \mathbb{N} \end{aligned}$$

$$x_t = x_{t-1} - \eta f'(x_{t-1})$$



# Optimization

## ❖ Square function



$$\begin{aligned} -100 \leq x \leq 100 \\ x \in \mathbb{N} \end{aligned}$$

$$x_t = x_{t-1} - \eta f'(x_{t-1})$$

$$x_0 = 70.0 \quad \eta = 0.1$$

$$f'(x_0) = 140.0$$

$$x_1 = x_0 - \eta f'(x_0) = 56.0$$

$$f'(x_1) = 112.0$$

$$x_2 = x_1 - \eta f'(x_1) = 44.8$$

$$f'(x_2) = 89.6$$

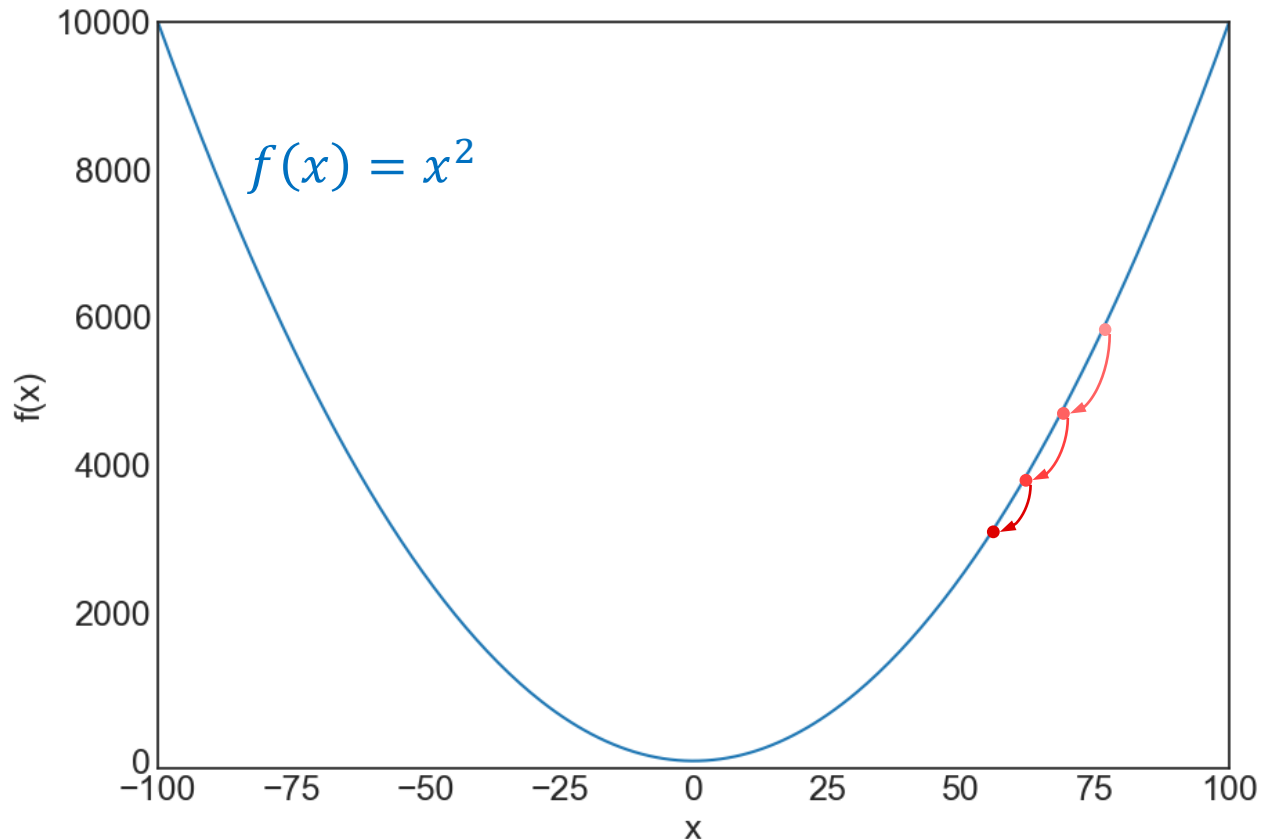
$$x_3 = x_2 - \eta f'(x_2) = 35.84$$

$$f'(x_3) = 71.68$$

$$x_4 = x_3 - \eta f'(x_3) = 28.672$$

# Optimization

## ❖ Square function



Keep doing

$$x_t = x_{t-1} - \eta f'(x_{t-1})$$

$$x_{10} = 6.012 \quad \eta = 0.1$$

$$f'(x_{10}) = 12.02$$

$$x_{11} = x_{10} - \eta f'(x_{10}) = 4.81$$

$$f'(x_{11}) = 9.62$$

$$x_{12} = x_{11} - \eta f'(x_{11}) = 3.84$$

$$f'(x_{12}) = 7.69$$

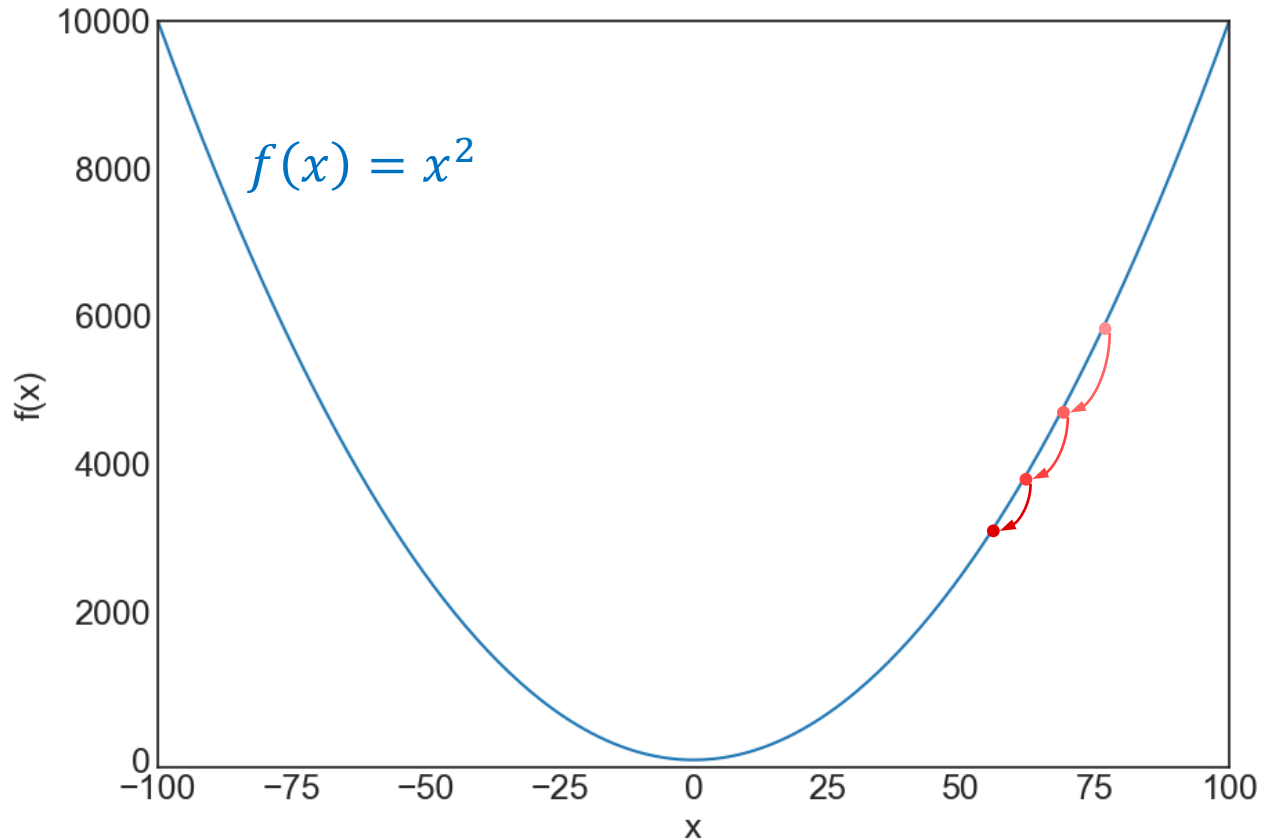
$$x_{13} = x_{12} - \eta f'(x_{12}) = 3.078$$

$$f'(x_{13}) = 6.15$$

$$x_{14} = x_{13} - \eta f'(x_{13}) = 2.46$$

# Optimization

## ❖ Square function



Keep doing

$$x_t = x_{t-1} - \eta f'(x_{t-1})$$

$$x_{30} = 0.069 \quad \eta = 0.1$$

$$f'(x_{30}) = 0.138$$

$$x_{31} = x_{30} - \eta f'(x_{30}) = 0.055$$

$$f'(x_{31}) = 0.11$$

$$x_{32} = x_{31} - \eta f'(x_{31}) = 0.044$$

$$f'(x_{32}) = 0.88$$

$$x_{33} = x_{32} - \eta f'(x_{32}) = 0.035$$

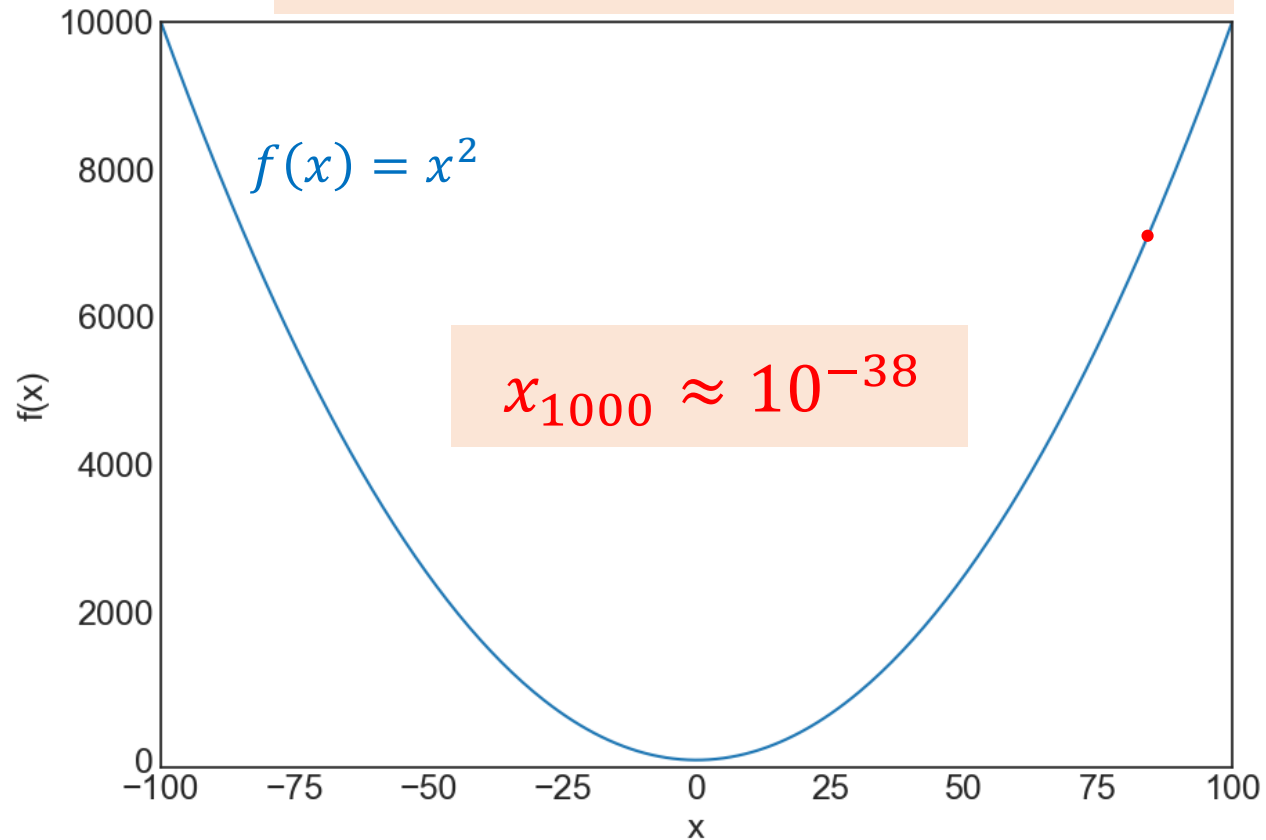
$$f'(x_{34}) = 0.071$$

$$x_{34} = x_{33} - \eta f'(x_{33}) = 0.028$$

# Optimization

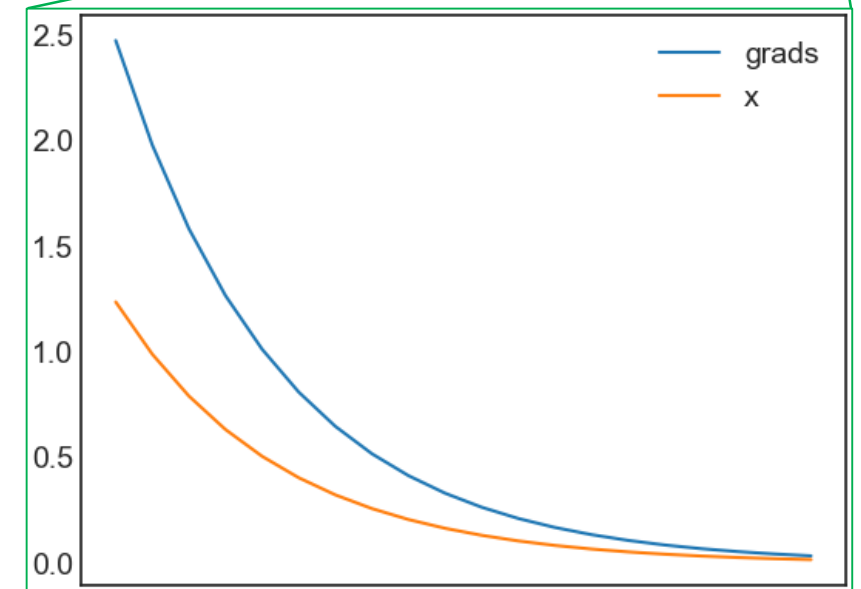
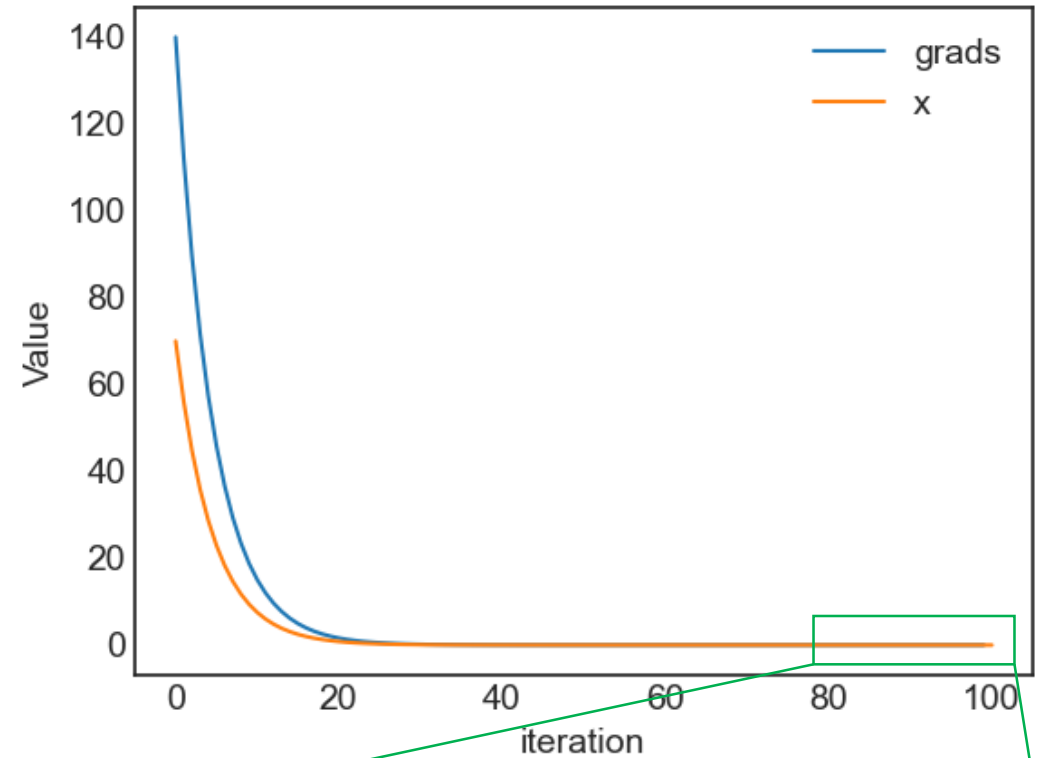
## ❖ Square function

$$x_t = x_{t-1} - \eta f'(x)$$



$$x_{1000} \approx 10^{-38}$$

Optimized successfully!



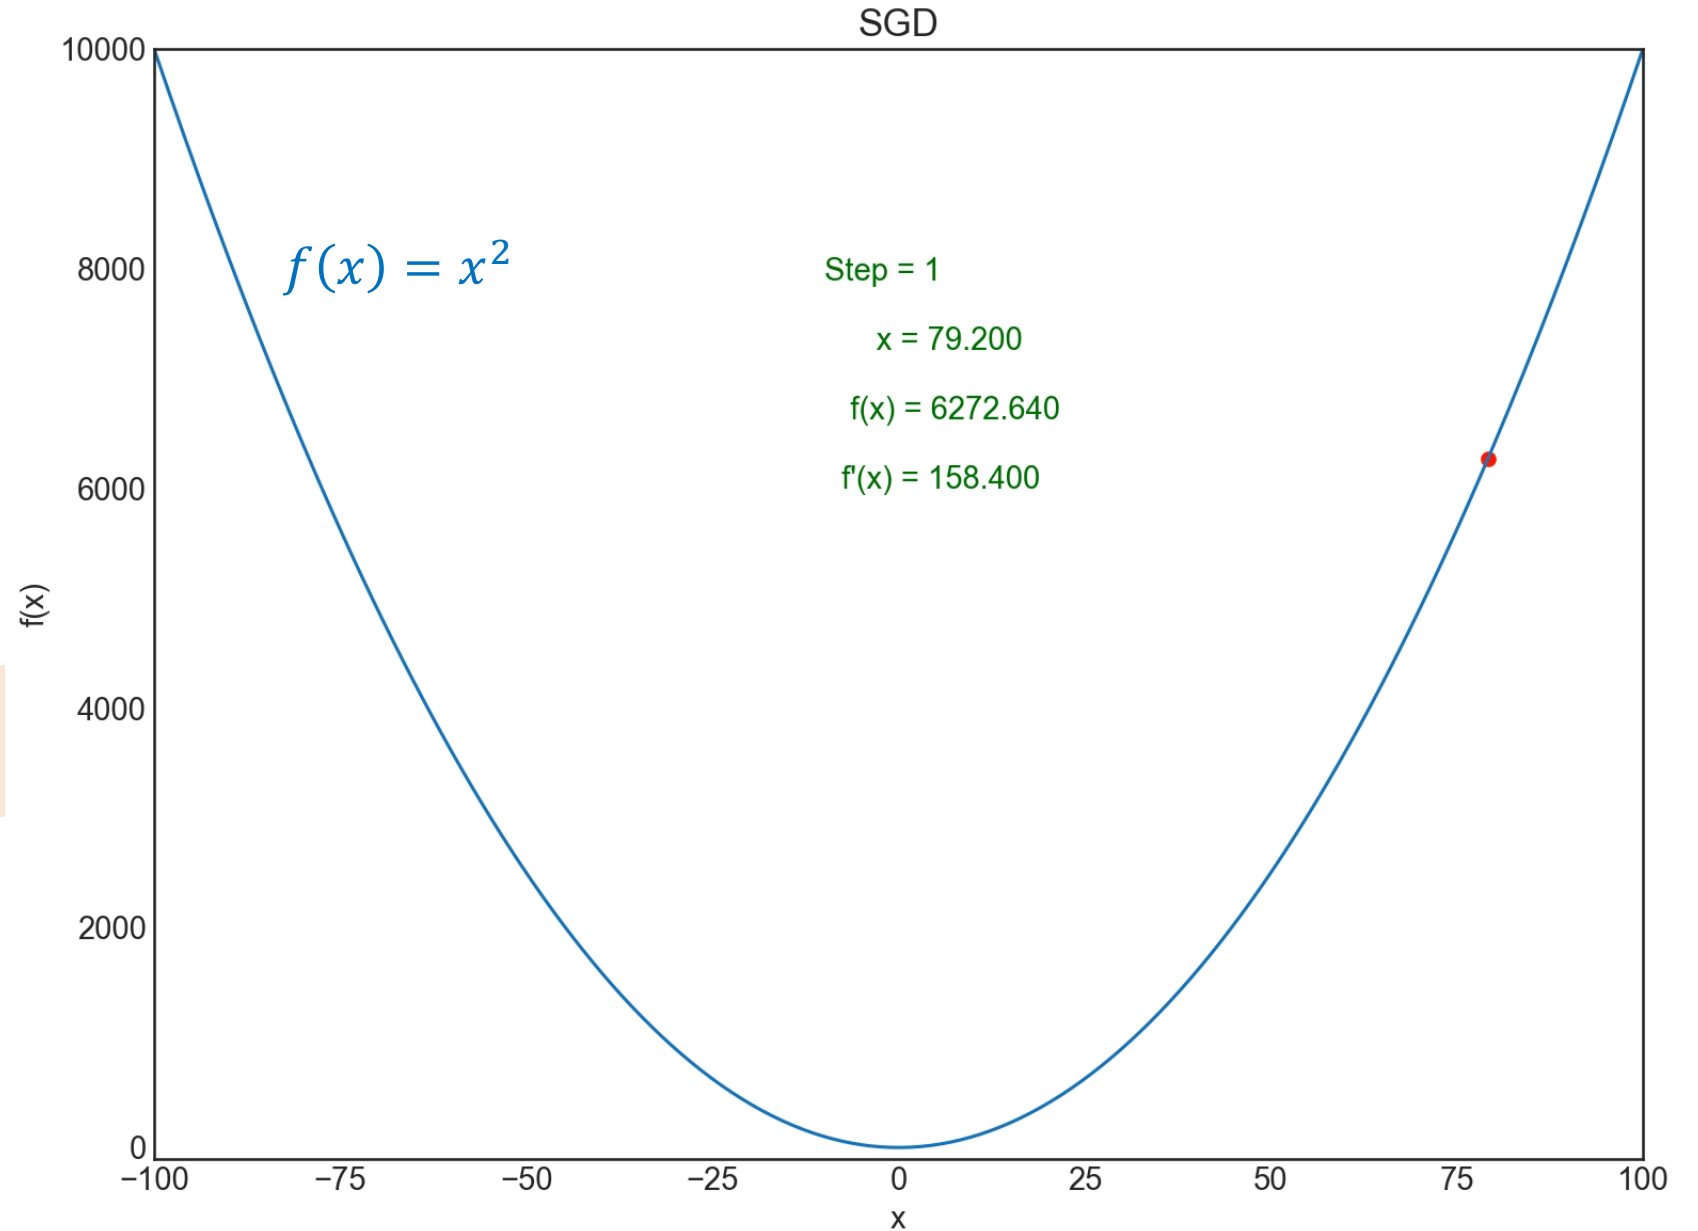
# Optimization

## ❖ Square function

$$x_0 = 99.0$$

$$\eta = 0.1$$

$$x_t = x_{t-1} - \eta f'(x)$$



# Optimization

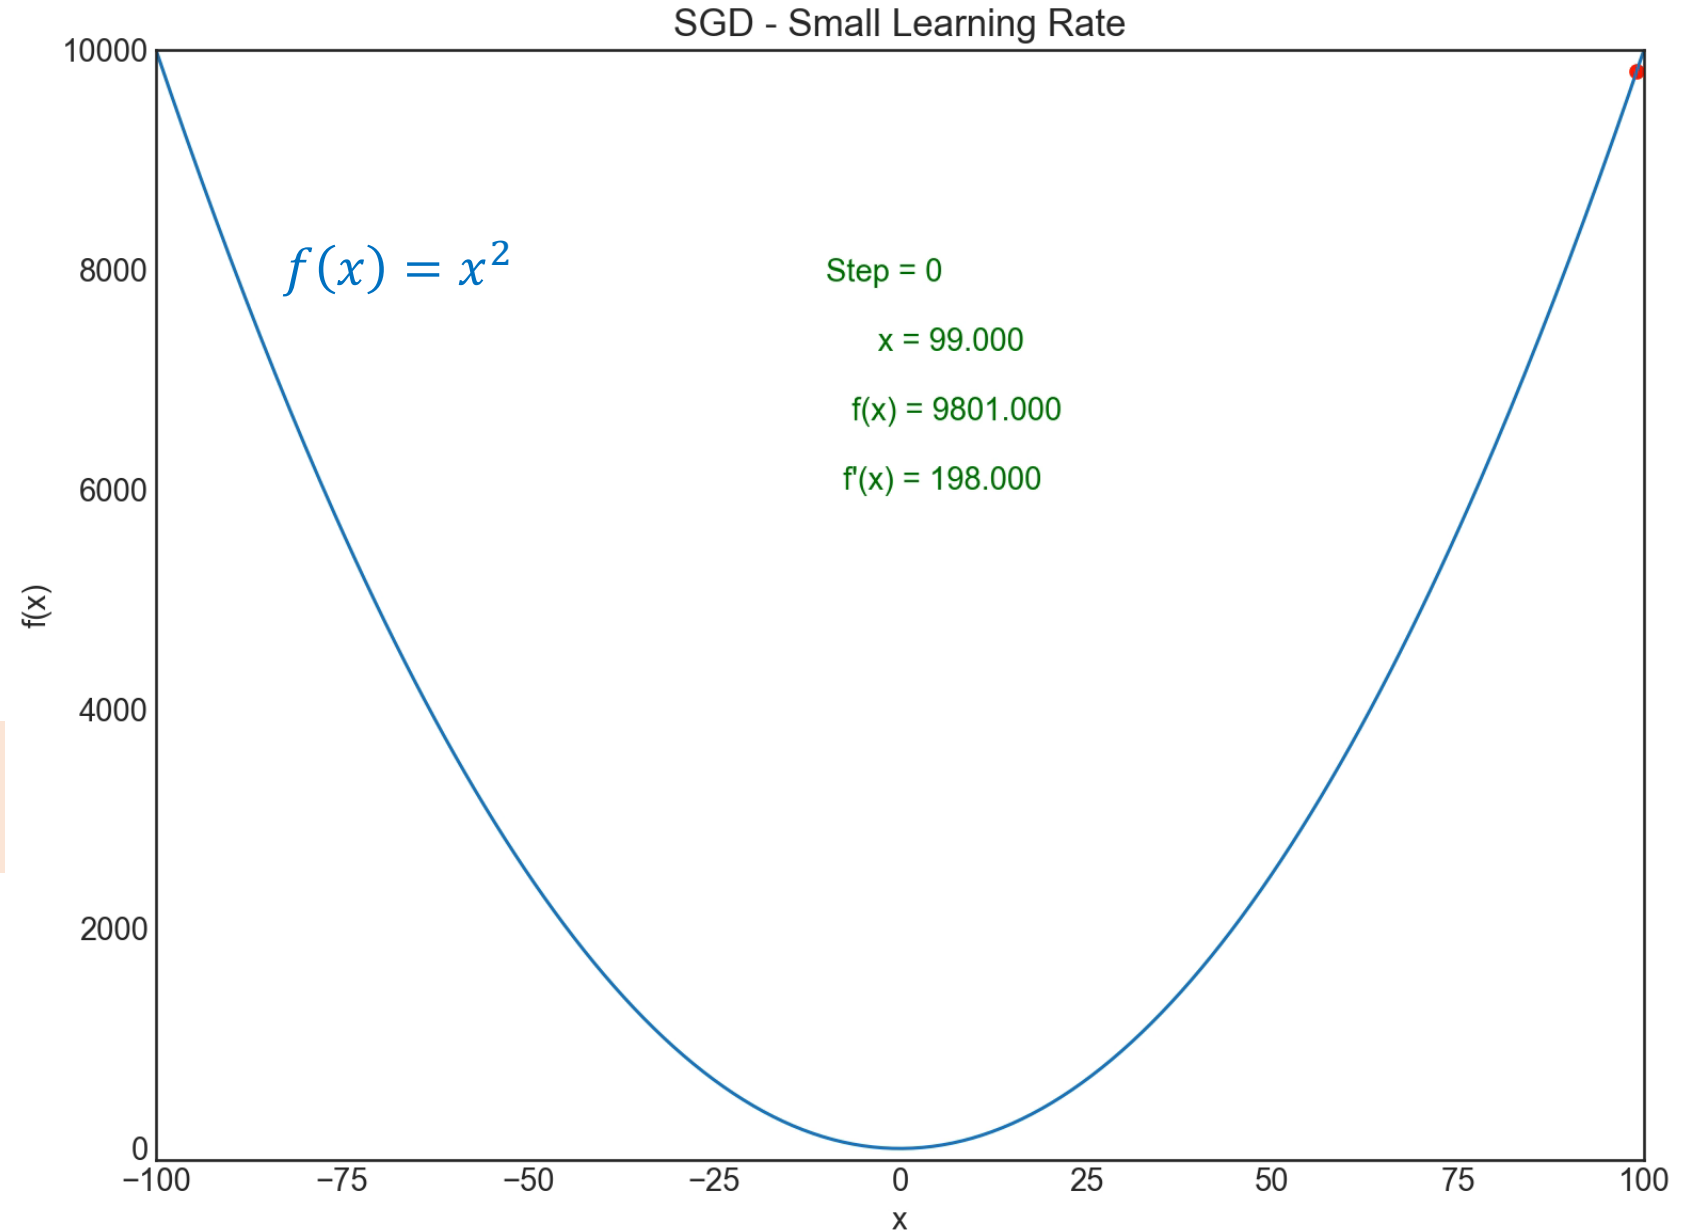
## ❖ Square function

### Discussion

$$x_0 = 99.0$$

$$\eta = 0.001$$

$$x_t = x_{t-1} - \eta f'(x)$$





# Optimization

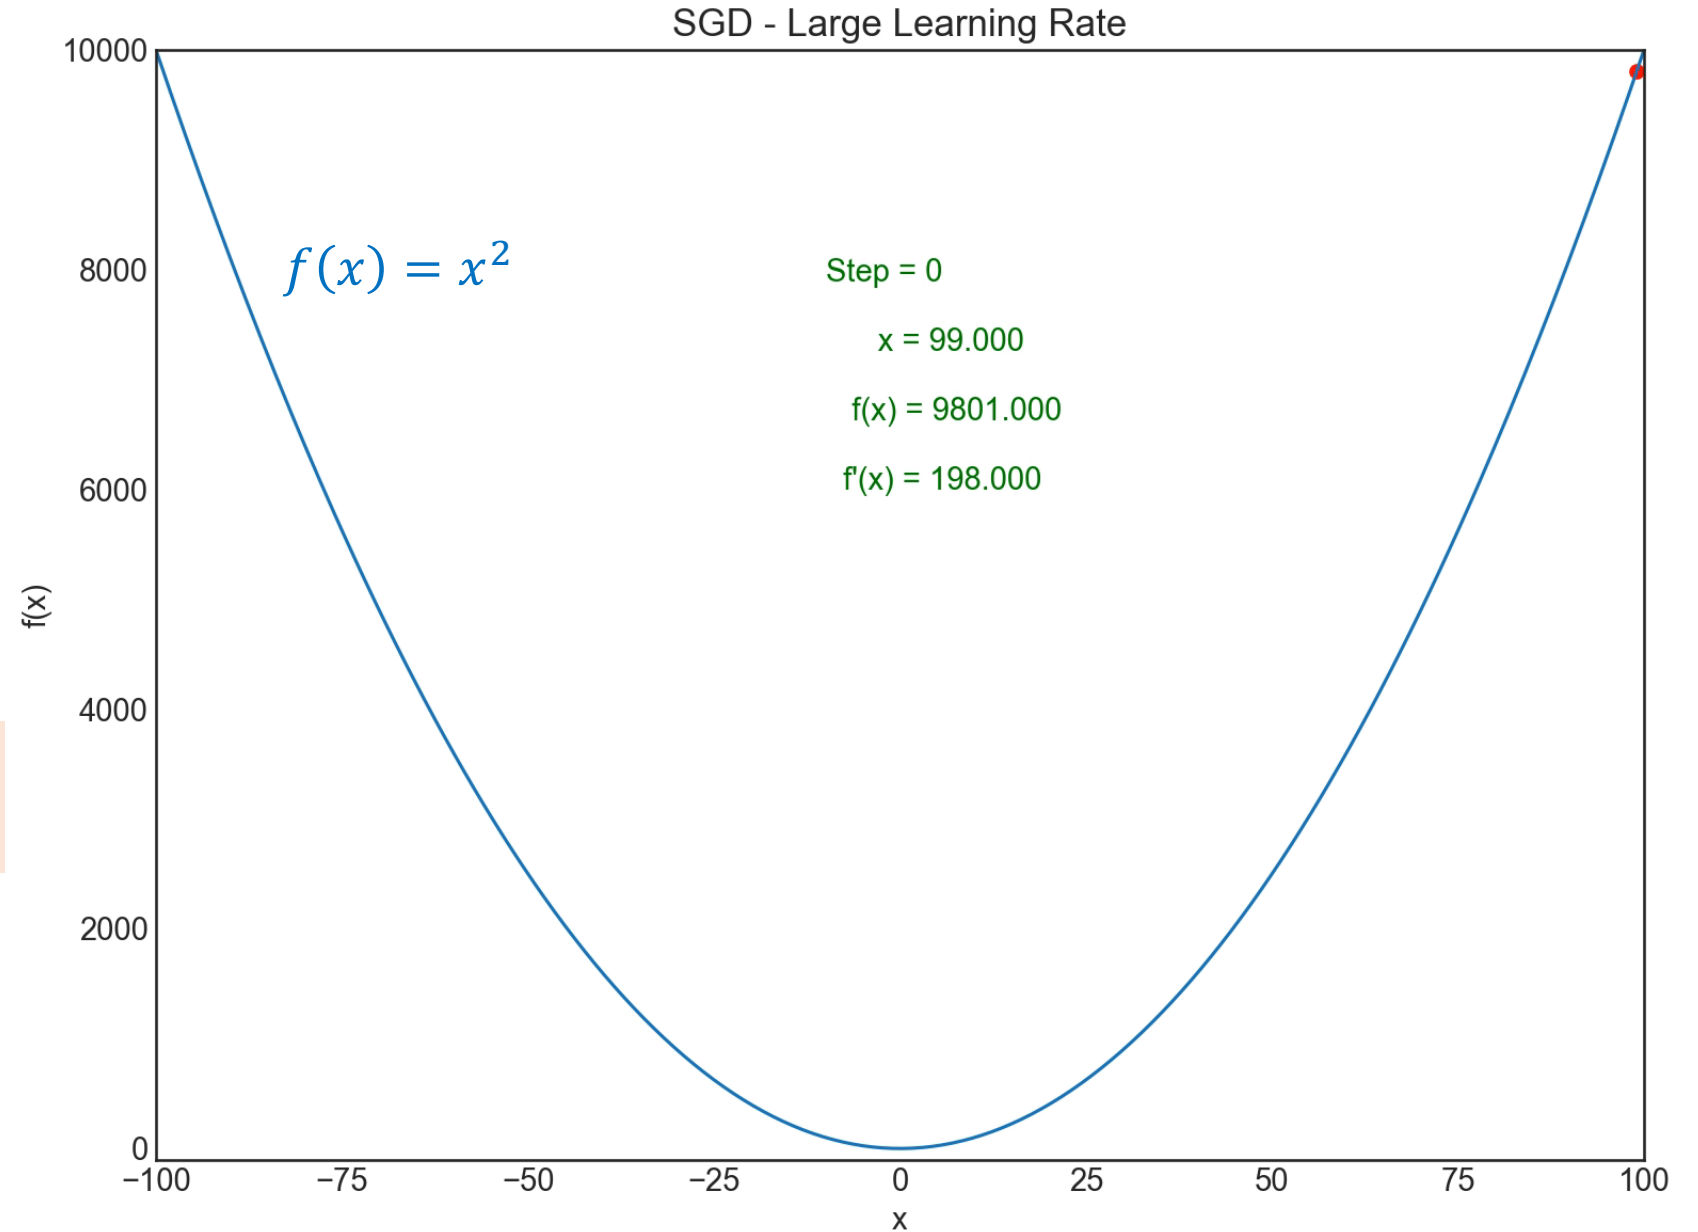
## ❖ Square function

### Discussion

$$x_0 = 99.0$$

$$\eta = 0.8$$

$$x_t = x_{t-1} - \eta f'(x)$$



# Optimization

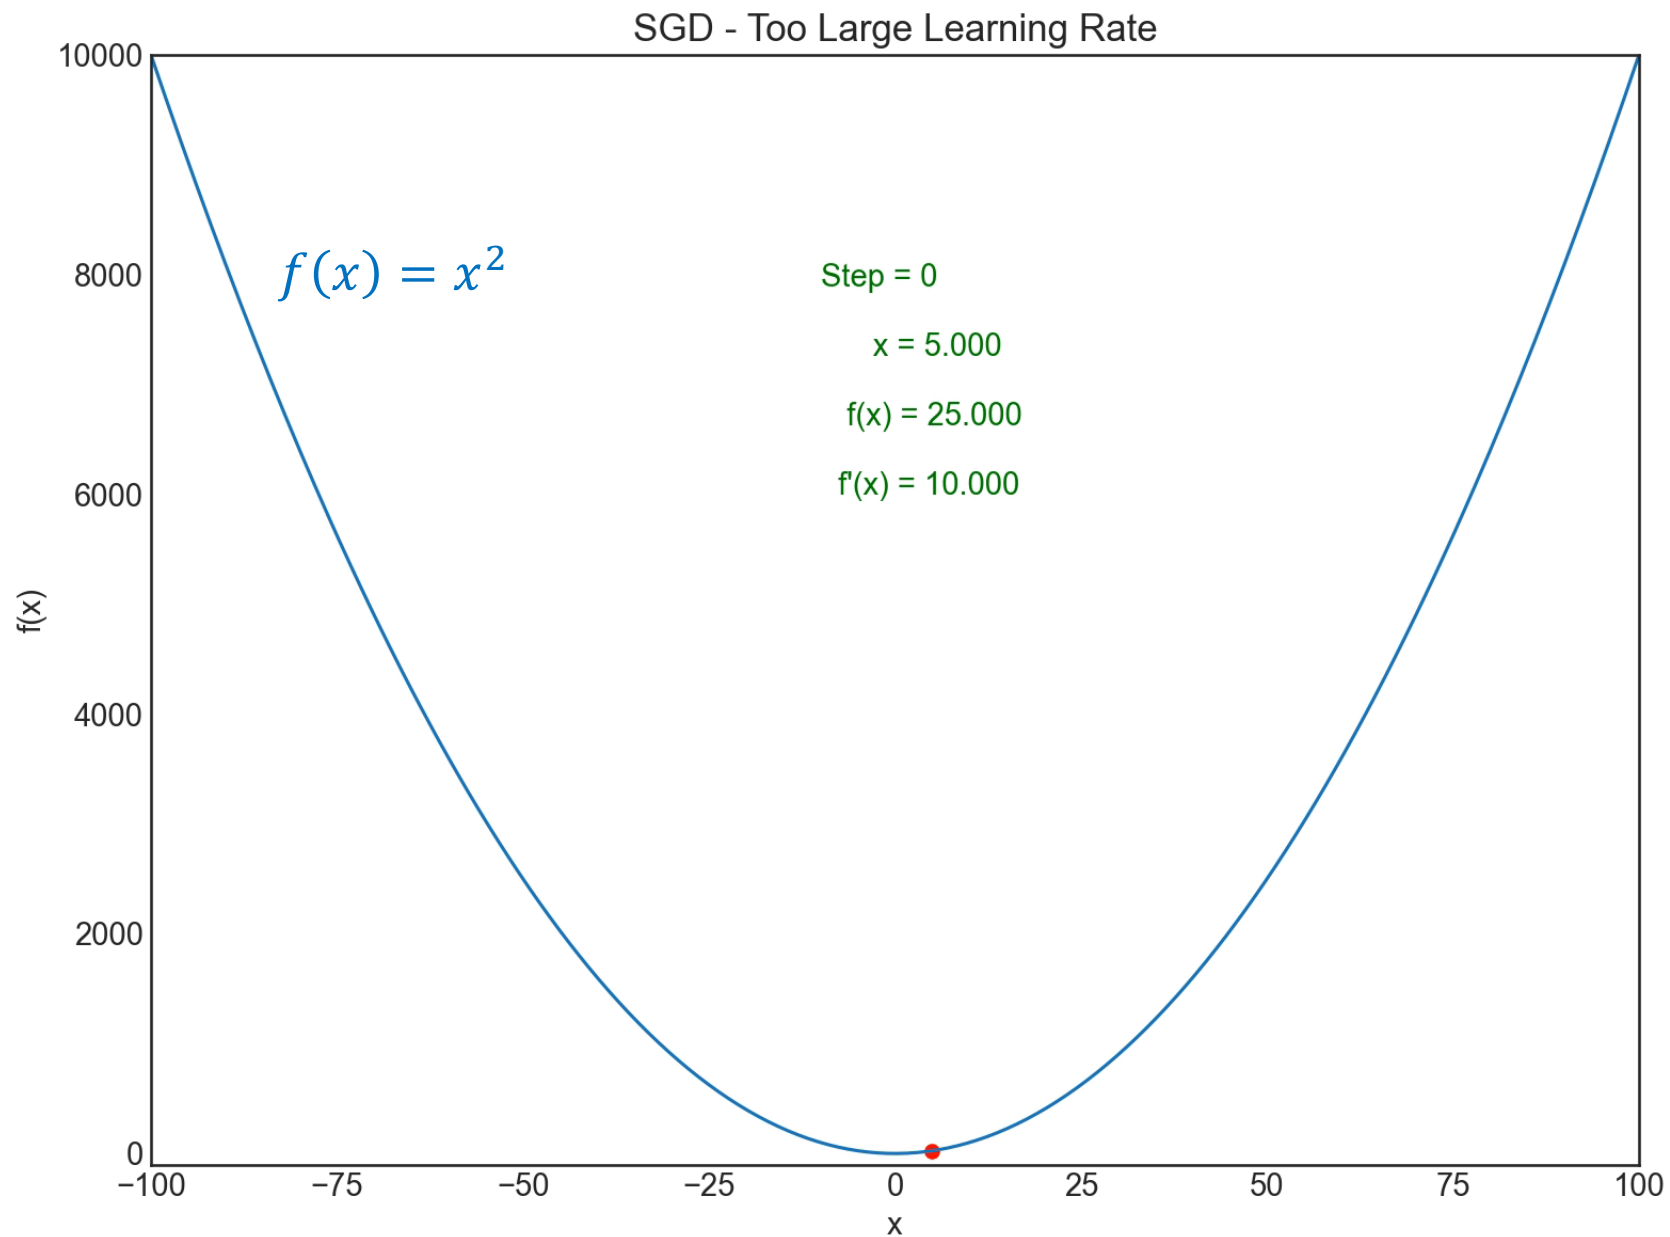
## ❖ Square function

### Discussion

$$x_0 = 99.0$$

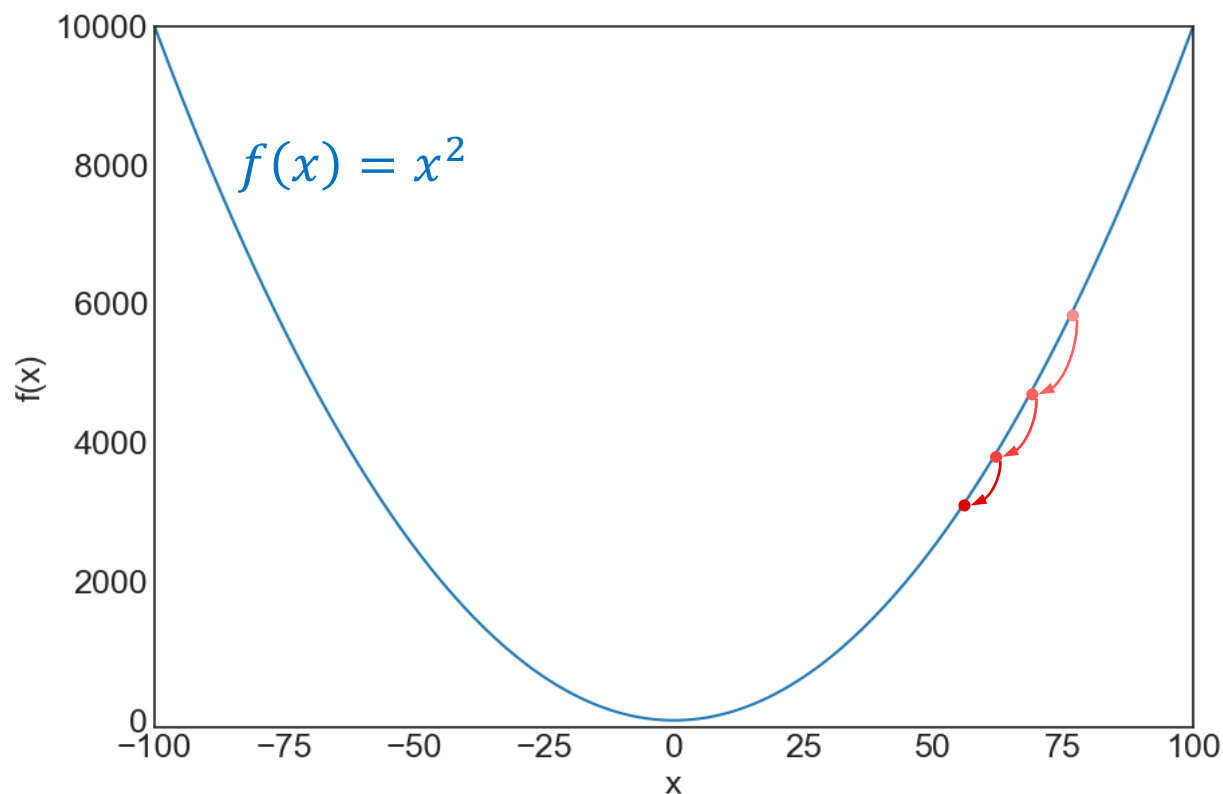
$$\eta = 1.1$$

$$x_t = x_{t-1} - \eta f'(x)$$



# Optimization

## ❖ Square function



- Given a function  $f(x)$ , find optimal  $x_{\text{opt}}$  so that  $f(x_{\text{opt}})$  is minimum
- After an update,  $f(x_{\text{new}}) \leq f(x_{\text{old}})$

$$x_t = x_{t-1} - \eta f'(x_{t-1})$$

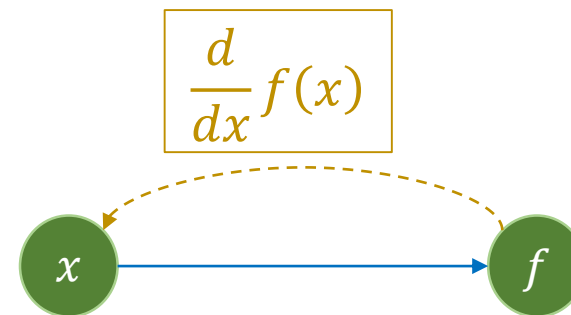
$$x_{30} = 0.069 \quad \eta = 0.1$$

$$f'(x_{30}) = 0.138$$

$$x_{31} = x_{30} - \eta f'(x_{30}) = 0.055$$

$$f'(x_{31}) = 0.11$$

$$x_{32} = x_{31} - \eta f'(x_{31}) = 0.044$$





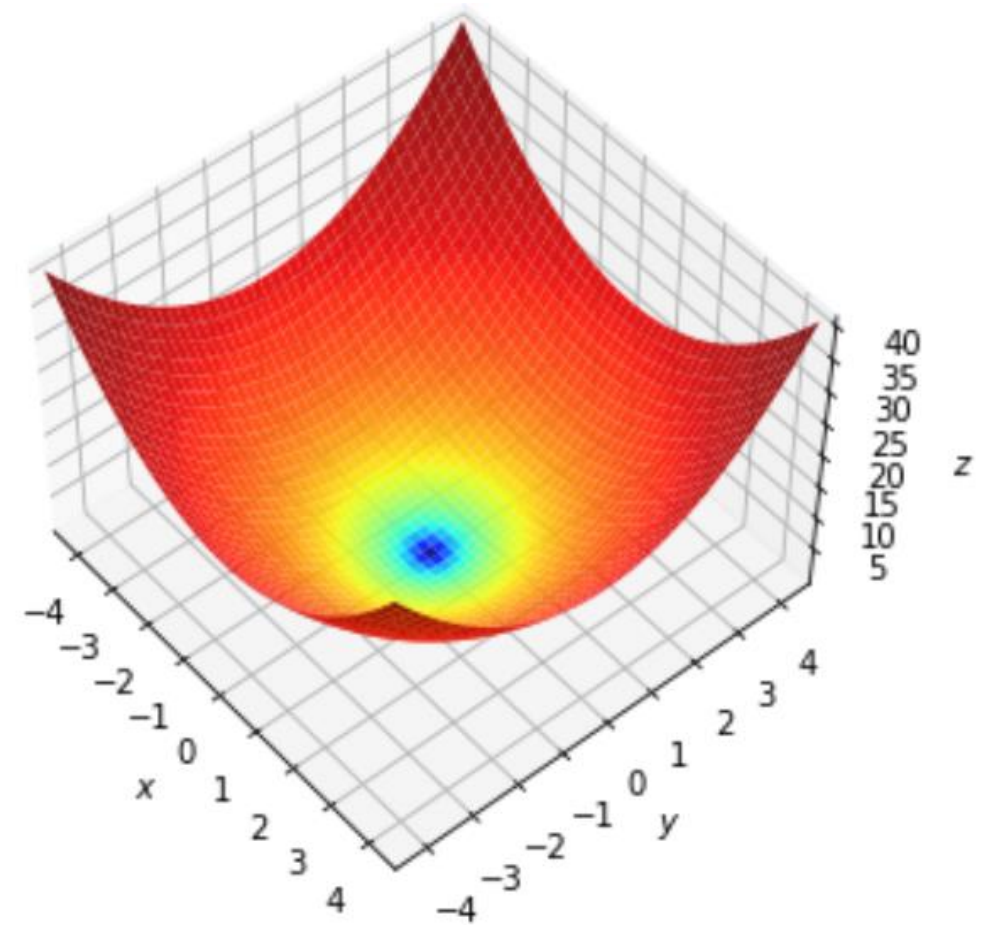
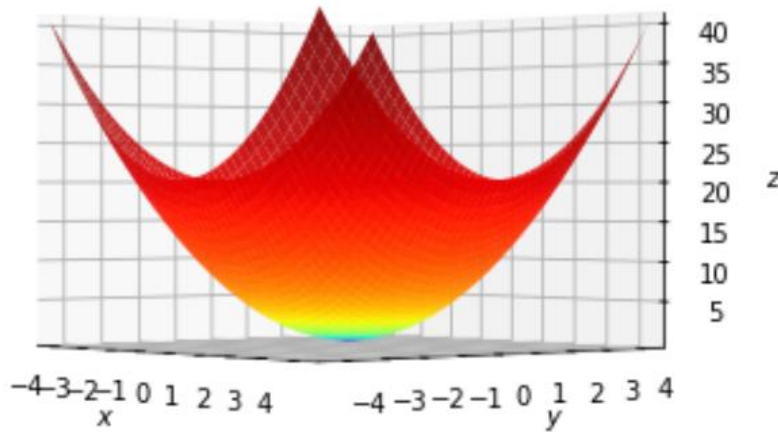
# Optimization

## ❖ Optimization: 2D function

$$f(x, y) = x^2 + y^2$$

$$-100 \leq x, y \leq 100$$

$$x, y \in \mathbb{N}$$



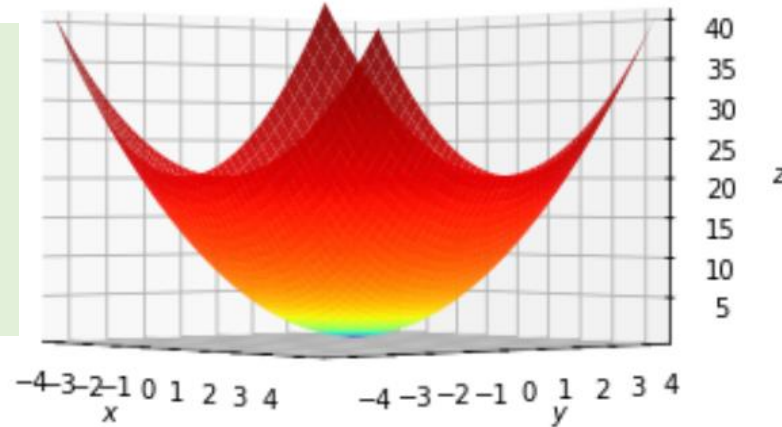
# Optimization

## ❖ Optimization: 2D function

$$f(x, y) = x^2 + y^2$$

$$-100 \leq x, y \leq 100$$

$$x, y \in \mathbb{N}$$



$$x = x - \eta \frac{\partial f(x, y)}{\partial x}$$

$$y = y - \eta \frac{\partial f(x, y)}{\partial y}$$

$$x_0 = 6.0$$

$$y_0 = 9.0$$

$$\eta = 0.1$$

$$\frac{\partial f(x_0, y_0)}{\partial x} = 12$$

$$x_1 = 4.8$$

$$\frac{\partial f(x_0, y_0)}{\partial y} = 18$$

$$y_1 = 7.2$$

$$\frac{\partial f(x_1, y_1)}{\partial x} = 9.6$$

$$x_2 = 3.84$$

$$\frac{\partial f(x_1, y_1)}{\partial y} = 14.4$$

$$y_2 = 5.75$$

$$\frac{\partial f(x_2, y_2)}{\partial x} = 7.68$$

$$x_3 = 3.07$$

$$\frac{\partial f(x_2, y_2)}{\partial y} = 11.51$$

$$y_3 = 4.608$$

$$\frac{\partial f(x_3, y_3)}{\partial x} = 6.14$$

$$x_4 = 2.45$$

$$\frac{\partial f(x_3, y_3)}{\partial y} = 9.21$$

$$y_4 = 3.68$$

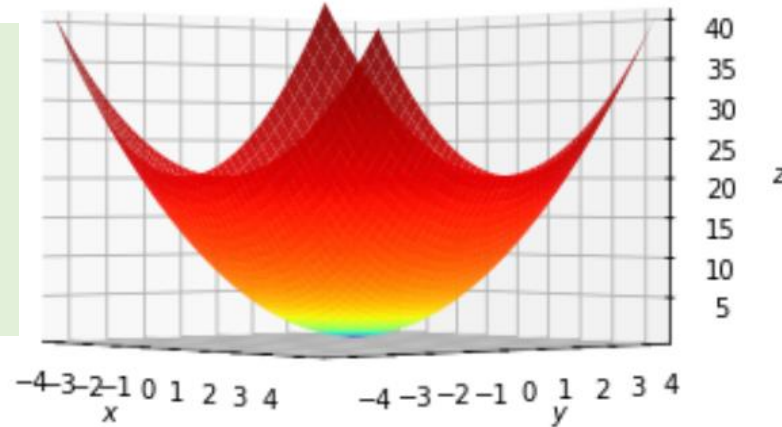
# Optimization

## ❖ Optimization: 2D function

$$f(x, y) = x^2 + y^2$$

$$-100 \leq x, y \leq 100$$

$$x, y \in \mathbb{N}$$



$$\frac{\partial f(x_2, y_2)}{\partial x} = 7.68 \quad \frac{\partial f(x_2, y_2)}{\partial y} = 11.51$$

$$x_3 = 3.07 \quad y_3 = 4.608$$

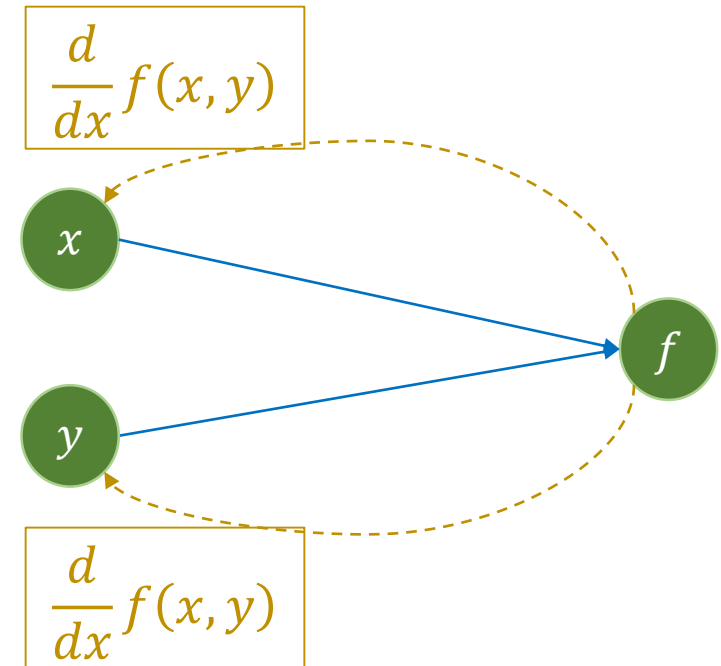


$$\frac{\partial f(x_3, y_3)}{\partial x} = 6.14 \quad \frac{\partial f(x_3, y_3)}{\partial y} = 9.21$$

$$x_4 = 2.45 \quad y_4 = 3.68$$

### Summary:

- Given a function  $f(x, y)$ , find optimal  $(x_{\text{opt}}, y_{\text{opt}})$  so that  $f(x_{\text{opt}}, y_{\text{opt}})$  is minimum
- After an update,  $f(x_{\text{new}}, y_{\text{new}}) \leq f(x_{\text{old}}, y_{\text{old}})$



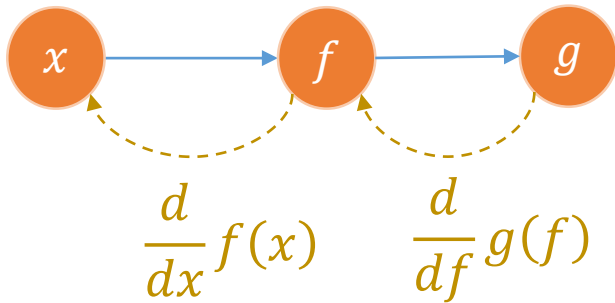
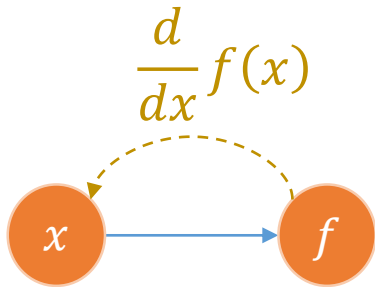
# Outline

- **Review on Optimization**
- **Partial Gradient and Chain Rules**
- **Finding a Line**
- **Linear Regression**

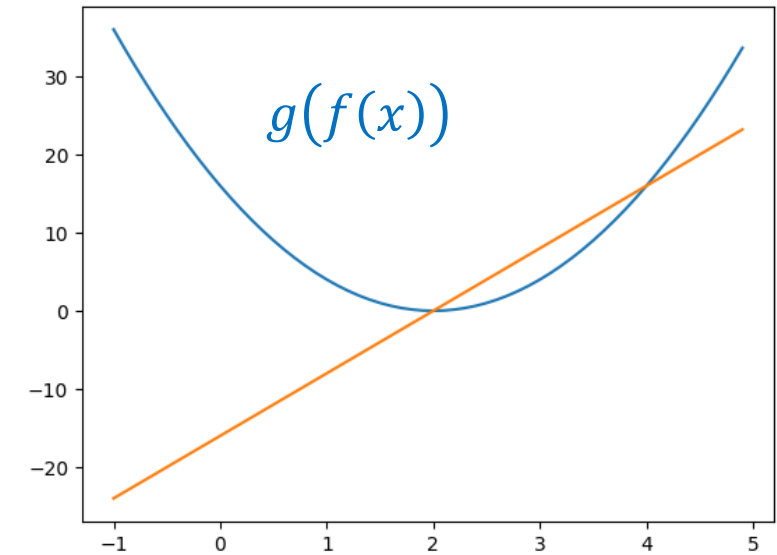
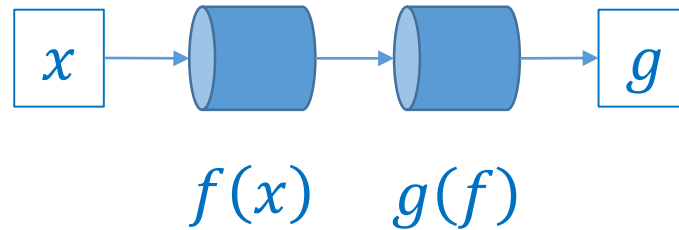
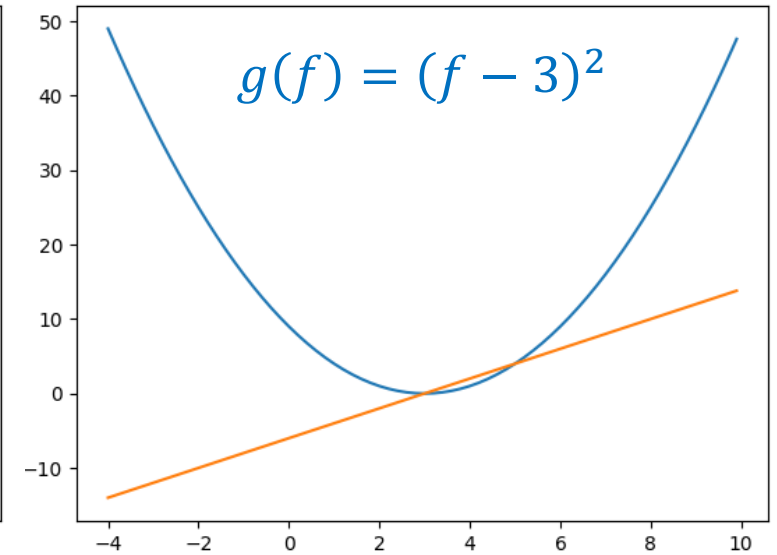
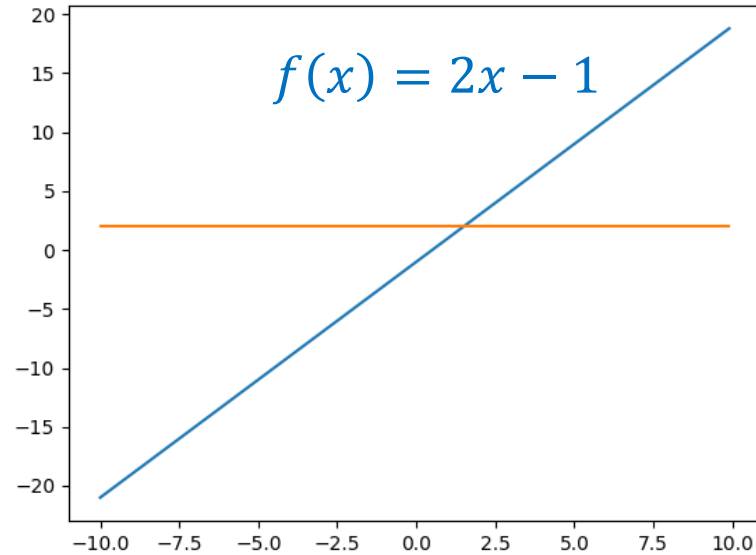


# Gradient-based Optimization

## ❖ For composite function

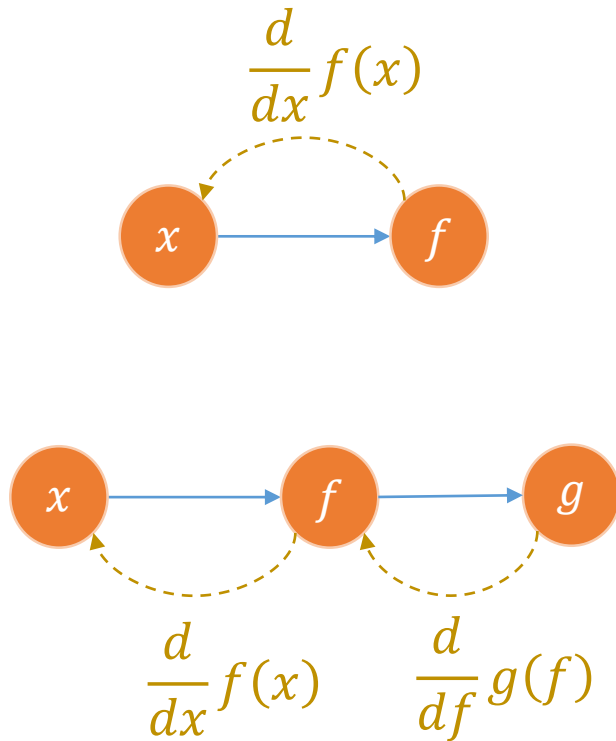


$$\frac{d}{dx} g(f(x)) = \left[ \frac{d}{df} g(f) \right] * \left[ \frac{d}{dx} f(x) \right]$$



# Gradient-based Optimization

## ❖ For composite function



$$\frac{d}{dx} g(f(x)) = \left[ \frac{d}{df} g(f) \right] * \left[ \frac{d}{dx} f(x) \right]$$



$$f(x) = 2x - 1$$

$$g(f) = (f - 3)^2$$

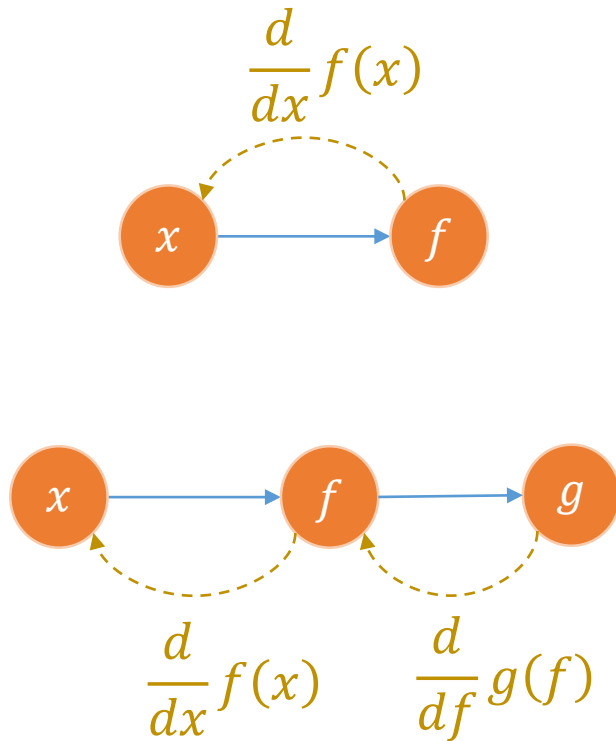
$$\begin{aligned} g(x) &= (2x - 1 - 3)^2 \\ &= (2x - 4)^2 \end{aligned}$$



$$\begin{aligned} g'(x) &= 4(2x - 4) \\ &= 8x - 16 \end{aligned}$$

# Gradient-based Optimization

## ❖ For composite function and chain rule



$$\frac{d}{dx} g(f(x)) = \left[ \frac{d}{df} g(f) \right] * \left[ \frac{d}{dx} f(x) \right]$$



$$f(x) = 2x - 1$$

$$g(f) = (f - 3)^2$$

$$f'(x) = 2$$

$$g'(f) = 2(f - 3)$$



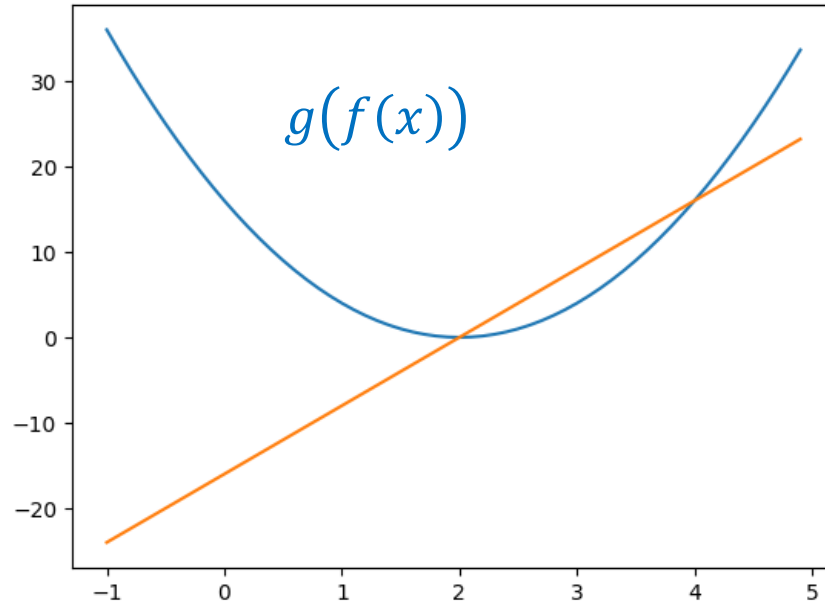
$$\begin{aligned} \frac{dg}{dx} &= \frac{dg}{df} \frac{df}{dx} \\ &= 2(f - 3)2 \\ &= 4(2x - 1 - 3) \\ &= 8x - 16 \end{aligned}$$

## Implementation

$$f(x) = 2x - 1$$

$$g(f) = (f - 3)^2$$

$$\frac{dg}{dx} = 8x - 16$$



```
1 def fx(x):  
2     return 2*x - 1  
3  
4 def gf(f):  
5     return (f-3)**2  
6  
7 def dg_dx(x):  
8     return 8*x - 16
```

```
1 import random  
2  
3 # parameters  
4 lr = 0.1  
5  
6 # initialize x  
7 x = 60  
8  
9 old_loss = gf(fx(x)) # Logging  
10 print(f'old_loss: {old_loss}')
```

```
11  
12 # compute derivative  
13 dg_dx_value = dg_dx(x)  
14  
15 # update  
16 x = x - lr*dg_dx_value  
17  
18 new_loss = gf(fx(x)) # Logging  
19 print(f'new_loss: {new_loss}')
```

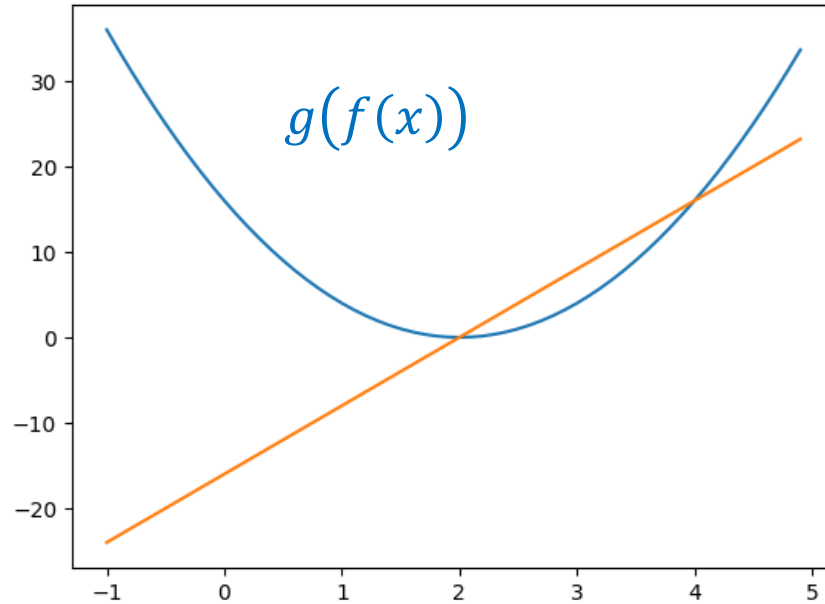
```
old_loss: 13456  
new_loss: 538.23999999999994
```

## Implementation

$$f(x) = 2x - 1$$

$$g(f) = (f - 3)^2$$

$$\frac{dg}{dx} = 8x - 16$$

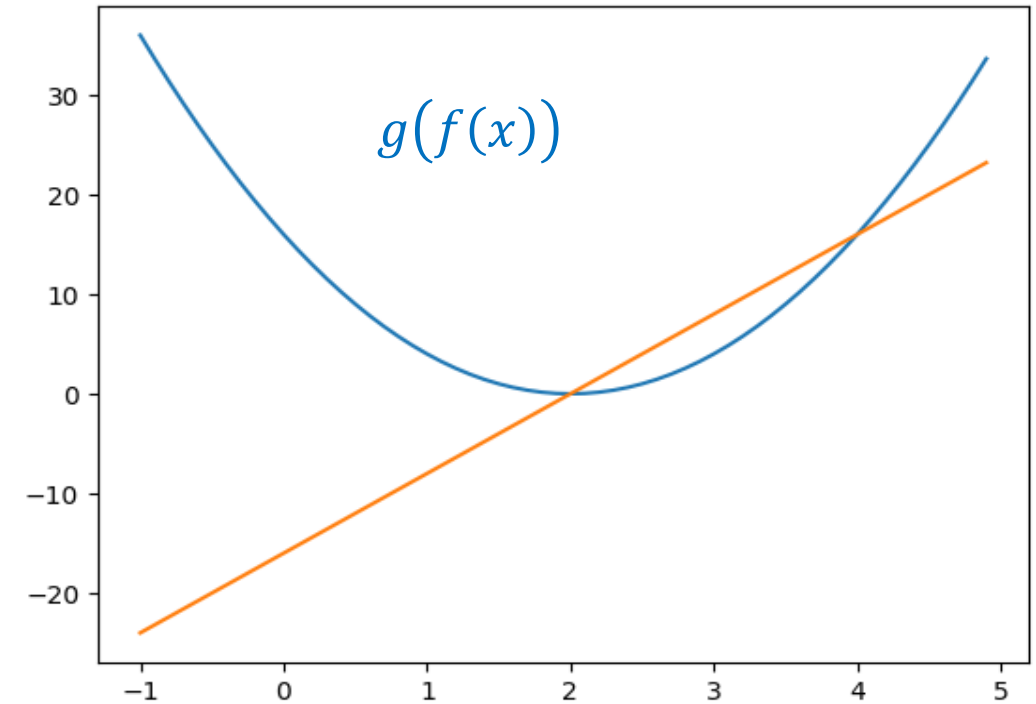
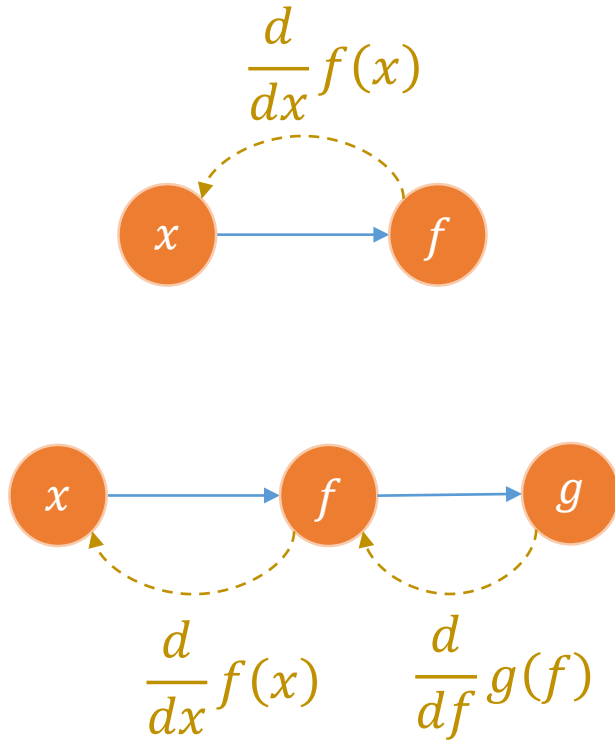


```
1 def fx(x):  
2     return 2*x - 1  
3  
4 def gf(f):  
5     return (f-3)**2  
6  
7 def dg_dx(x):  
8     return 8*x - 16
```

```
1 import random  
2  
3 # parameters  
4 num_steps = 5  
5 lr = 0.1  
6  
7 # set x randomly  
8 x = random.randint(-100, 100)  
9  
10 for _ in range(num_steps):  
11     # logging  
12     loss = gf(fx(x))  
13  
14     # compute derivative  
15     dg_dx_value = dg_dx(x)  
16  
17     # update  
18     x = x - lr*dg_dx_value
```

# Gradient-based Optimization

## ❖ For composite function



- Given a function  $g(f(x))$ , find optimal  $x_{\text{opt}}$  so that  $f(x_{\text{opt}})$  is minimum
- After an update,  $g(f(x_{\text{new}})) \leq g(f(x_{\text{old}}))$

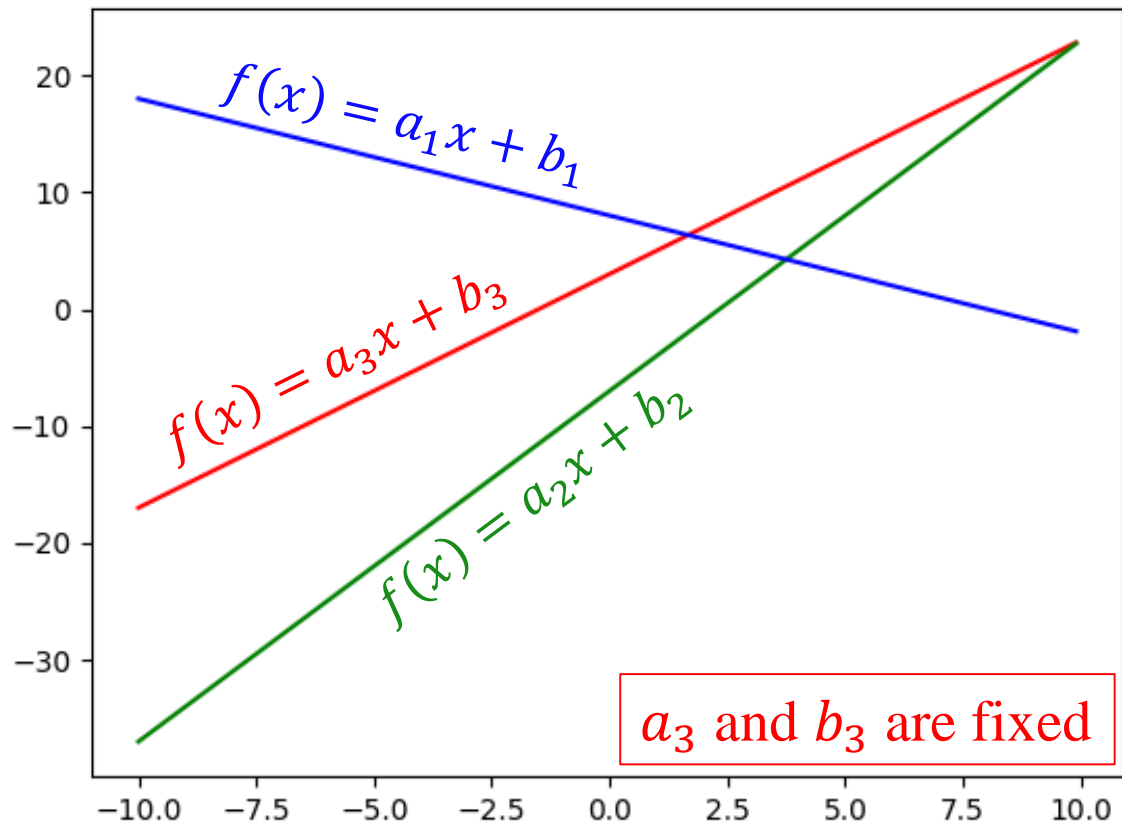
$$\frac{d}{dx} g(f(x)) = \left[ \frac{d}{df} g(f) \right] * \left[ \frac{d}{dx} f(x) \right]$$



# Gradient-based Optimization

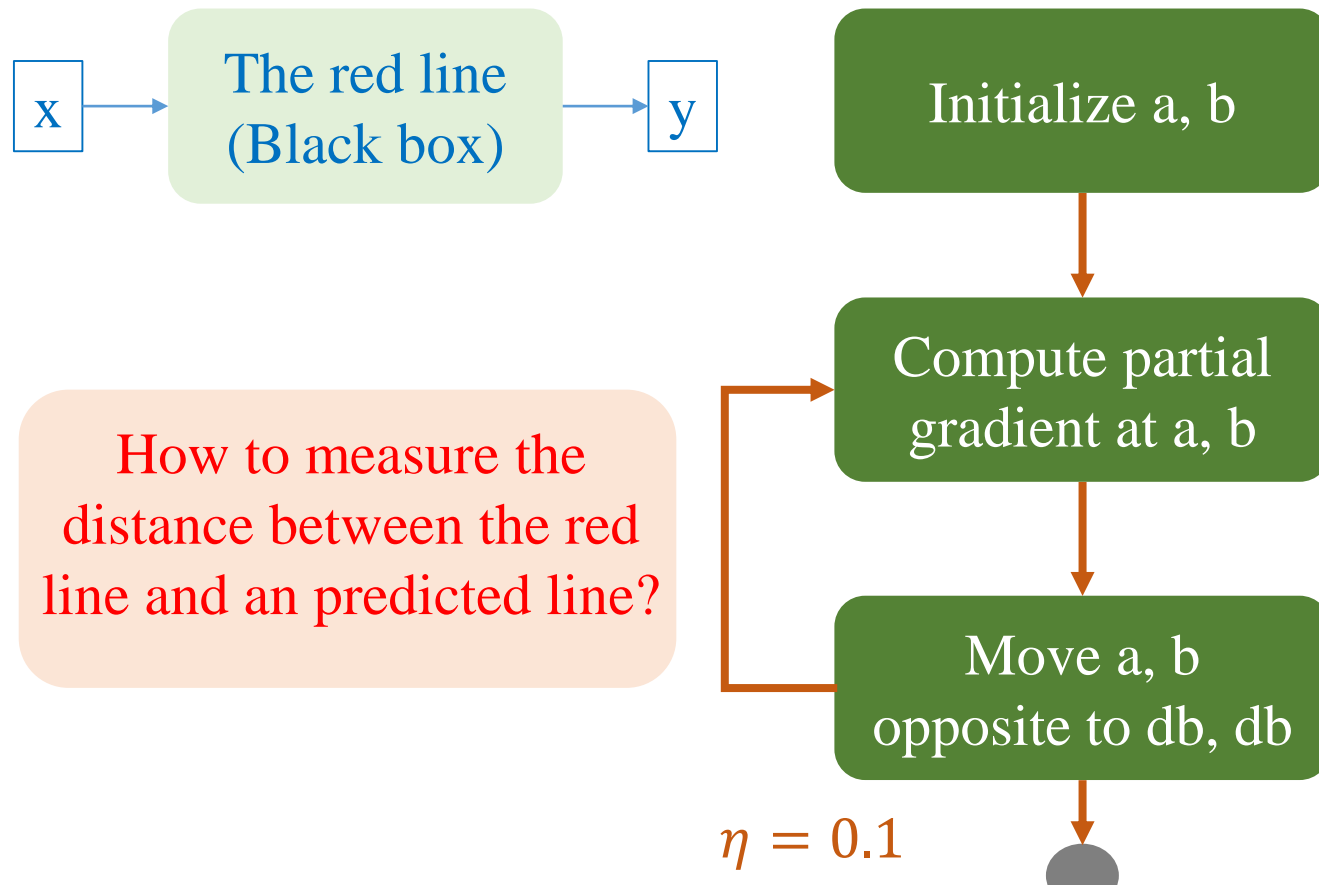
## ❖ Another context

$$f(x) = ax + b$$



## ❖ Constraints

$$x_t = x_{t-1} - \eta f'(x_{t-1})$$

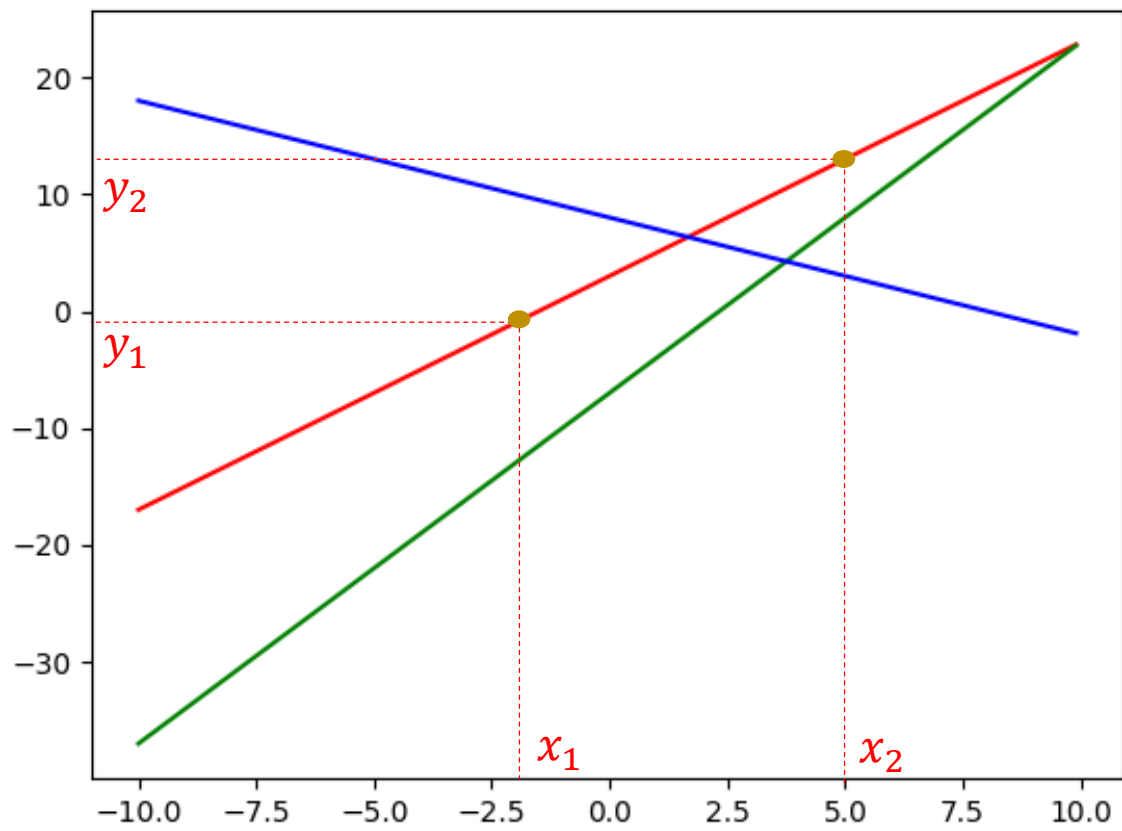




# Gradient-based Optimization

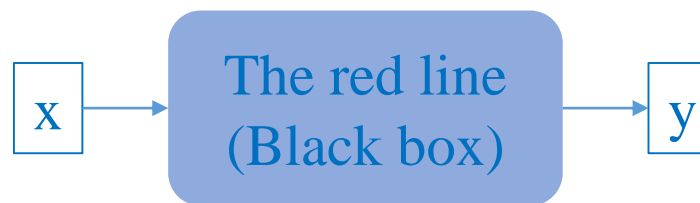
## ❖ Another context

$$f(x) = ax + b$$



## ❖ Constraints

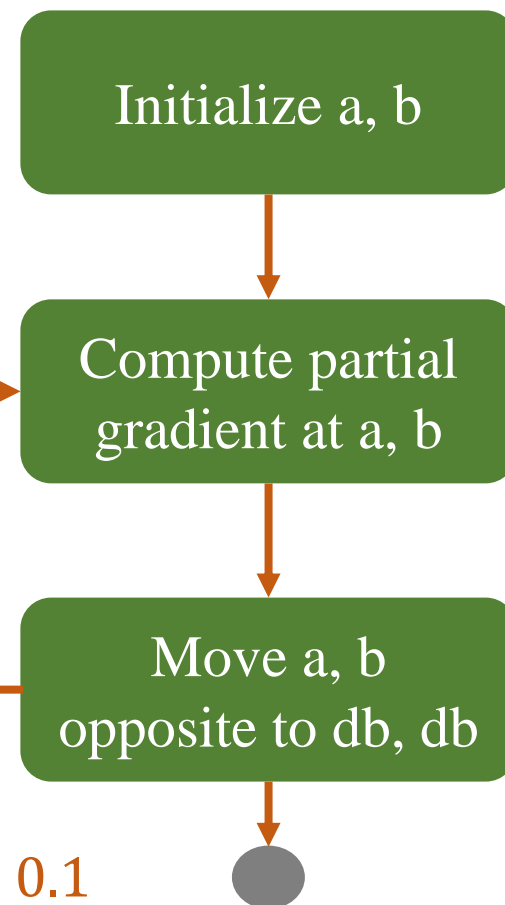
$$x_t = x_{t-1} - \eta f'(x_{t-1})$$



$$(x_1 = -2, y_1 = -1)$$

$$(x_2 = 5, y_2 = 13)$$

$$g(f) = (f - y_i)^2$$

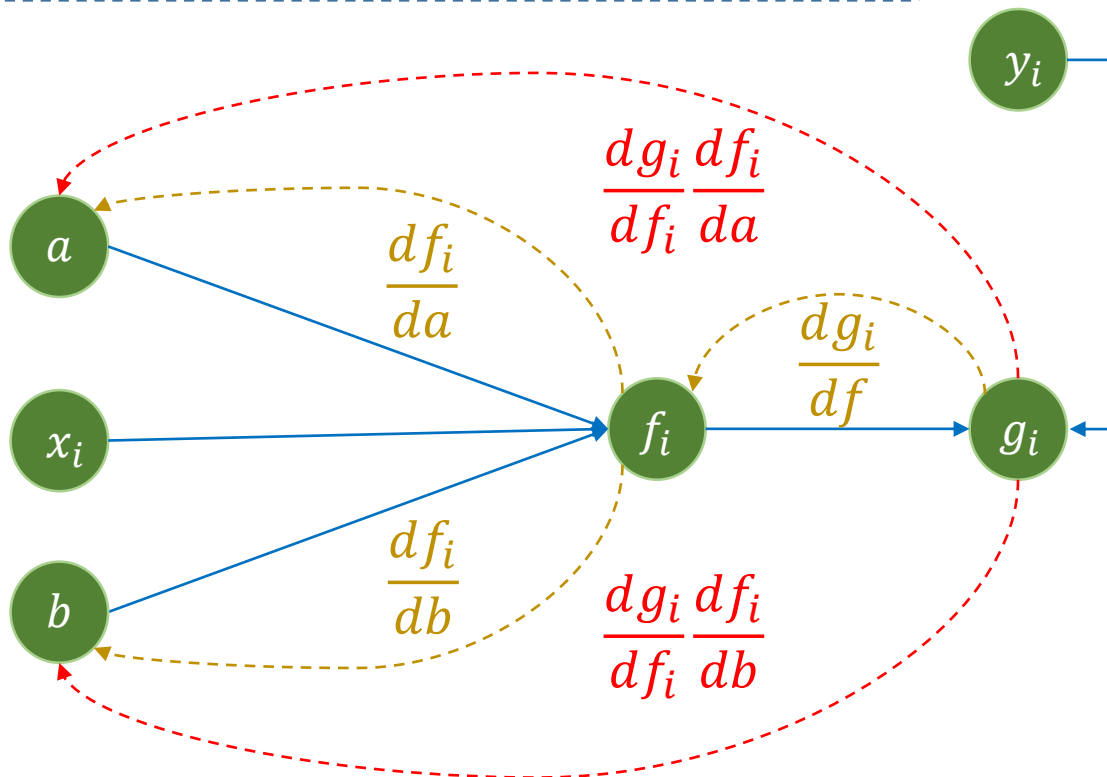


# Gradient-based Optimization

## ❖ Equations for partial gradients

$$f(x_i) = ax_i + b \quad (x_1=1, y_1=5)$$

$$g(f) = (f - y_i)^2 \quad (x_2=2, y_2=7)$$



$$\frac{df}{da} = x$$

$$\frac{df}{db} = 1$$

$$\frac{dg}{df} = 2(f - y)$$

$$\frac{dg}{da} = \frac{dg}{df} \frac{df}{da} = 2x(f - y)$$

$$\frac{dg}{db} = \frac{dg}{df} \frac{df}{db} = 2(f - y)$$

During looking for optimal  $a$  and  $b$ , at a given time,  $a$  and  $b$  have concrete values

## ❖ Optimization for a composite function

Find a and b so that  $g(f(x))$  is minimum

$$f(x_i) = ax_i + b \quad (x_1=1, y_1=5)$$

$$g(f) = (f - y_i)^2 \quad (x_2=2, y_2=7)$$

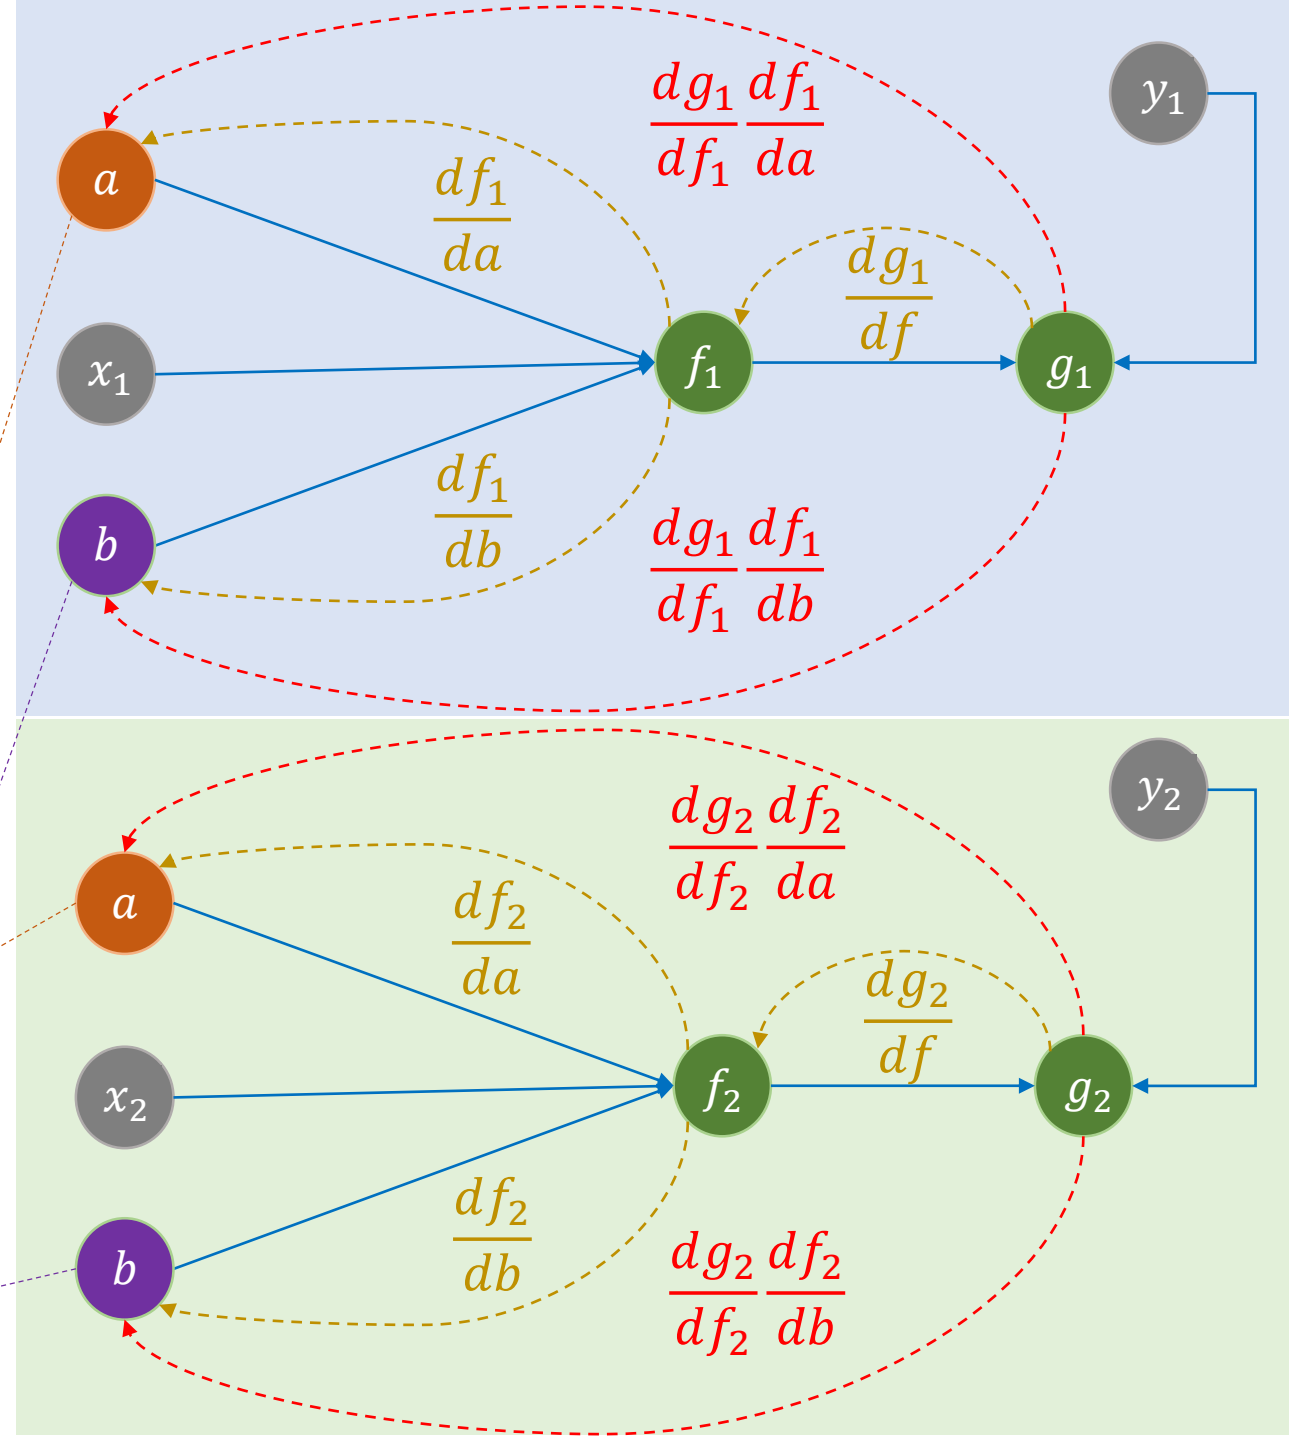
Partial derivative functions

$$\frac{dg}{da} = \frac{dg}{df} \frac{df}{da} = 2x(f - y)$$

$$\frac{dg}{db} = \frac{dg}{df} \frac{df}{db} = 2(f - y)$$

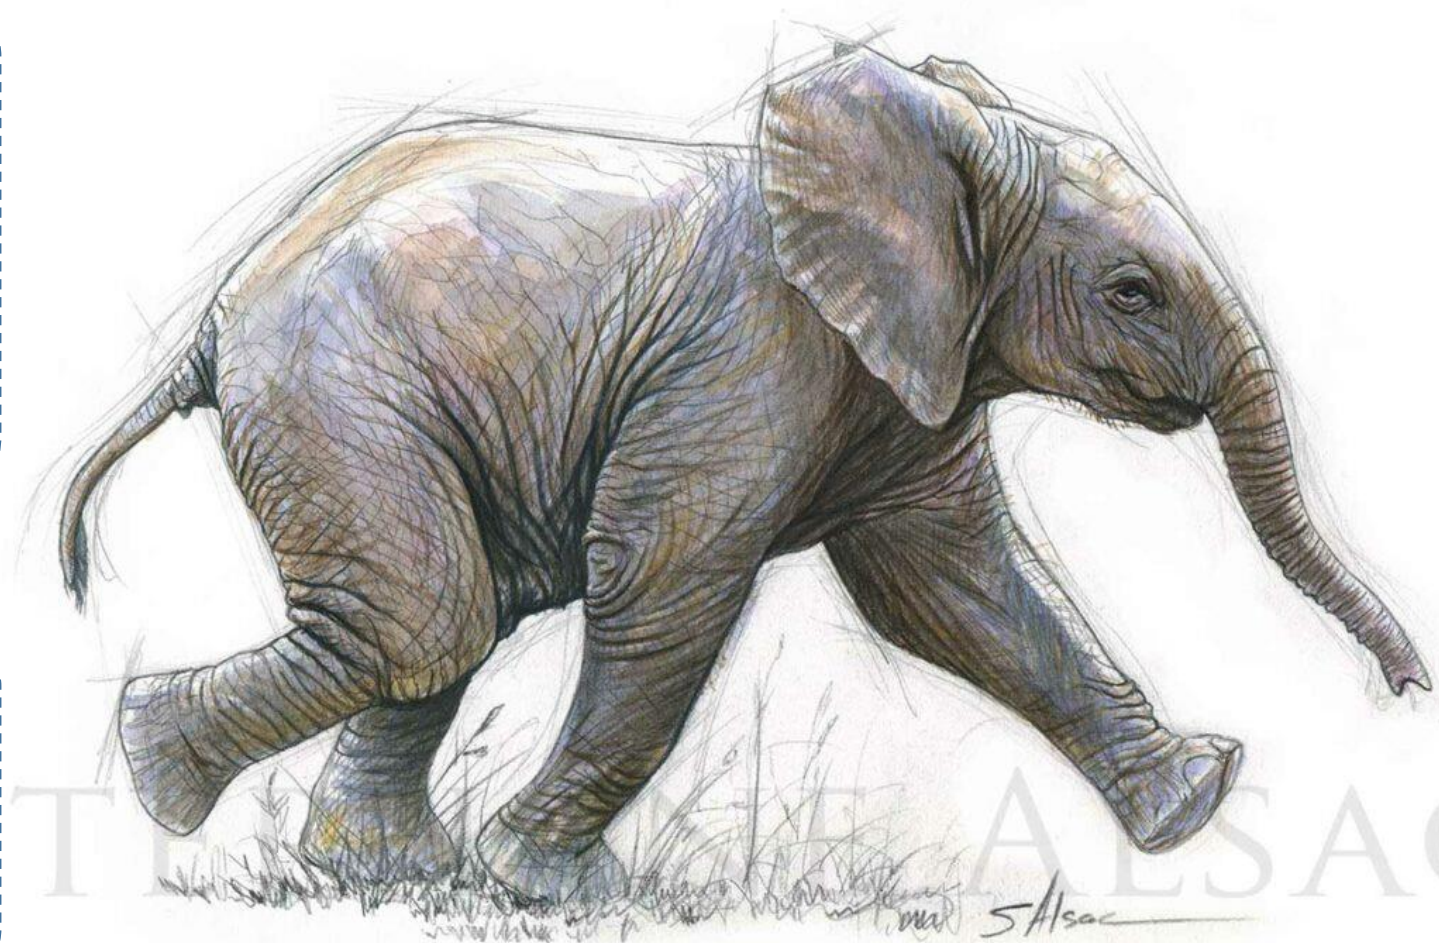
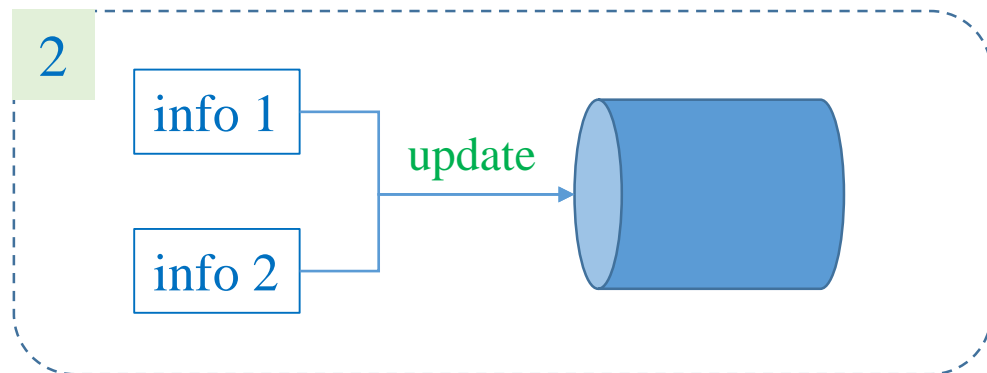
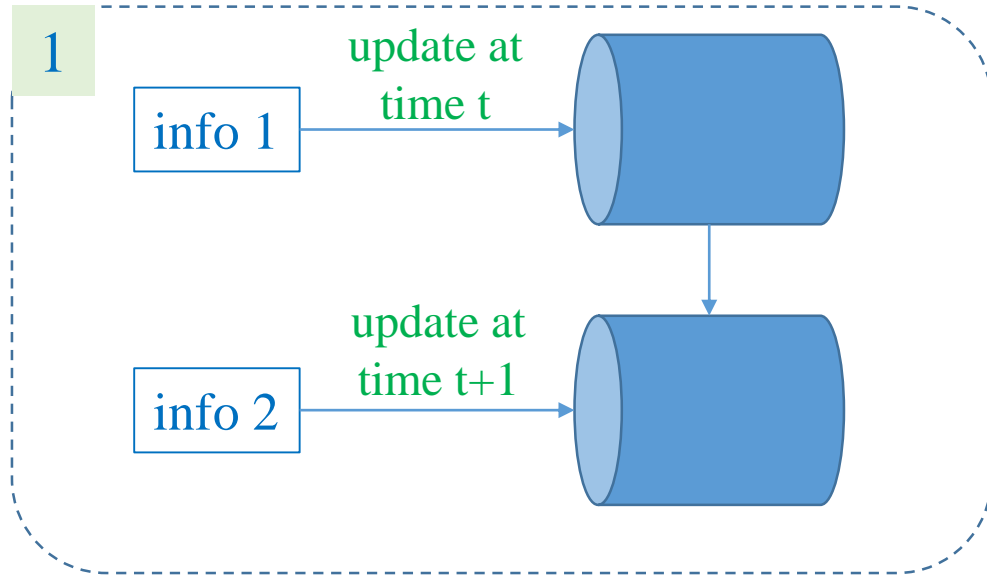
$$\sum_i \frac{dg_i}{da} = \frac{dg_1}{df_1} \frac{df_1}{da} + \frac{dg_2}{df_2} \frac{df_2}{da}$$

$$\sum_i \frac{dg_i}{db} = \frac{dg_1}{df_1} \frac{df_1}{db} + \frac{dg_2}{df_2} \frac{df_2}{db}$$



# Gradient-based Optimization

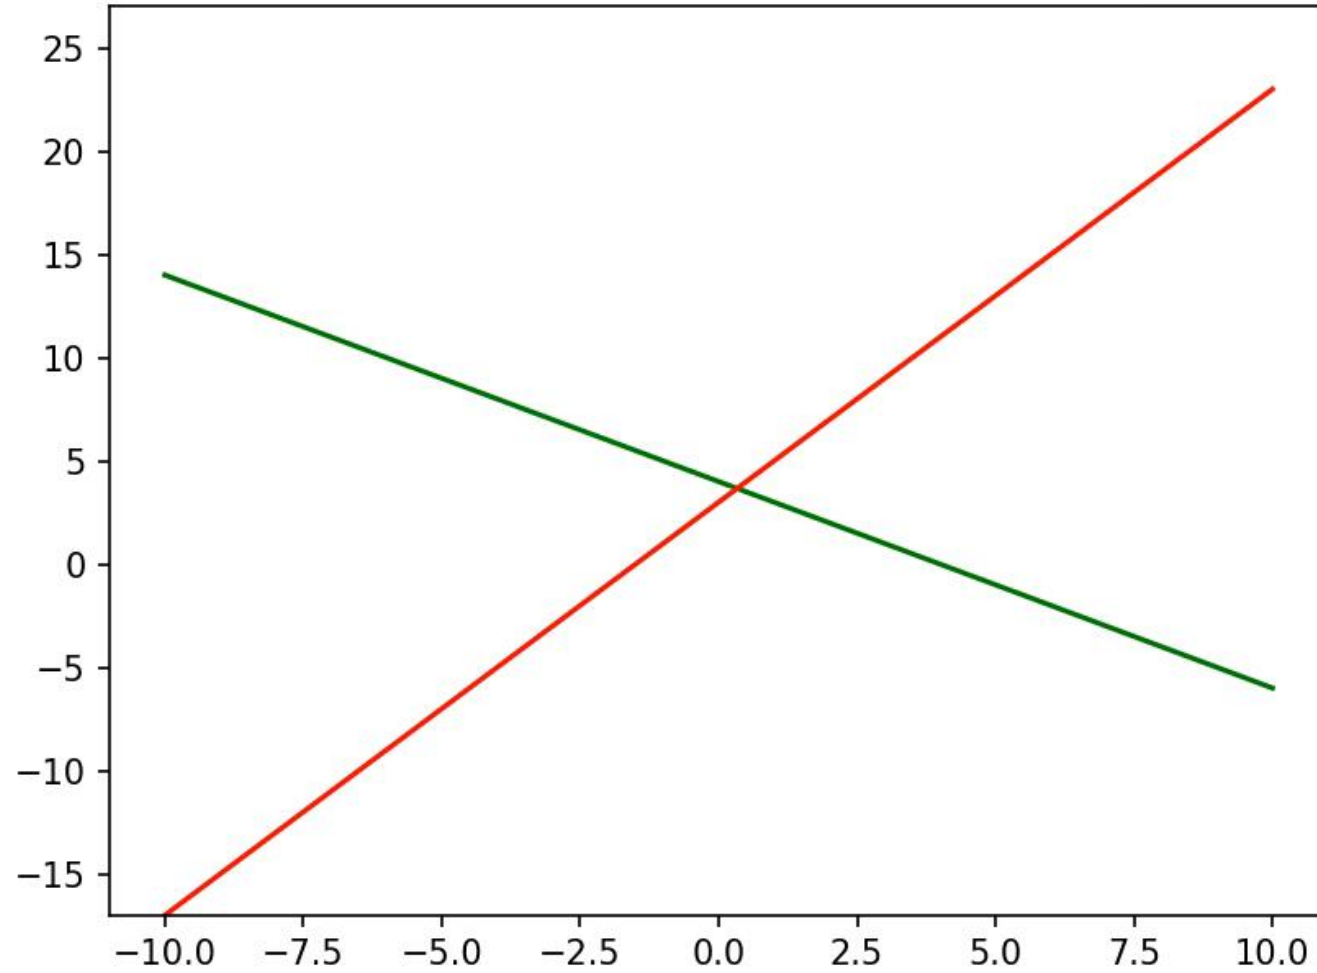
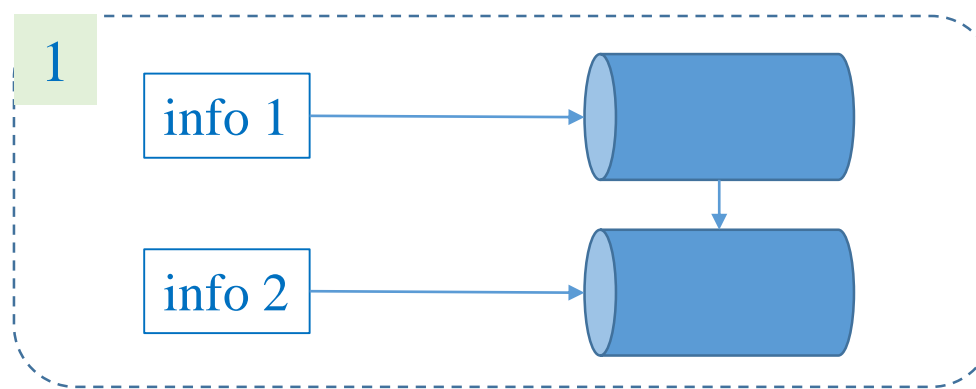
## ❖ How to use gradient information



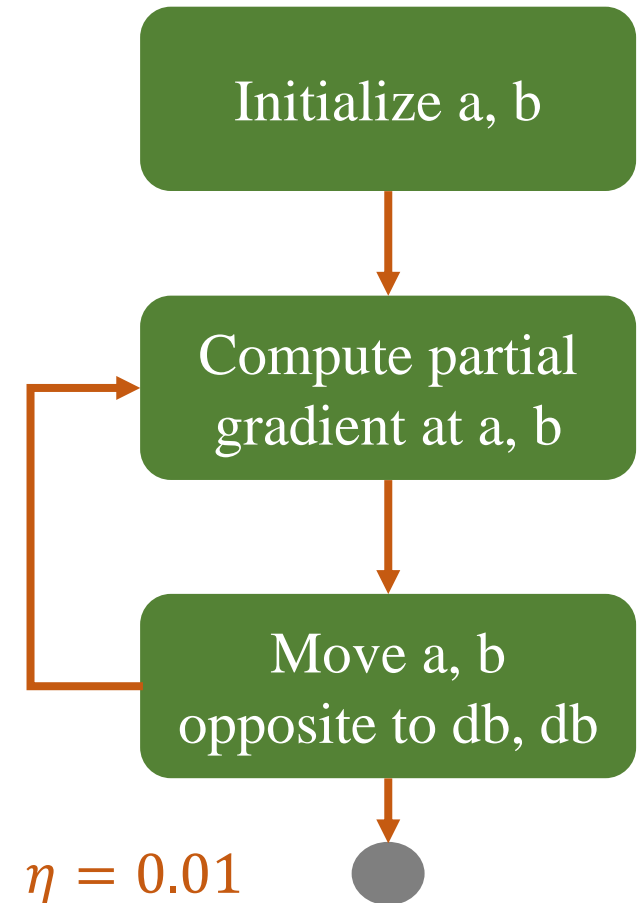
```
1 import random
2
3 # predict function
4 def predict_func(x, a, b):
5     return a*x + b
6
7 # parameters
8 num_steps = 100
9 lr = 0.01
10
11 # given data
12 data = [[-2, -1],
13         [5, 13]]
14
15 # 1. set a, b randomly
16 a = random.random()*10.0 - 5.0
17 b = random.random()*10.0 - 5.0
```

```
1 for i in range(num_steps):
2     for sample in data:
3         x_value, y_value = sample
4
5         # compute predicted_y
6         predicted_y = predict_func(x_value, a, b)
7
8         # compute g
9         g_value = (predicted_y - y_value)**2
10
11         # compute partial gradients for a and b
12         dg_da = 2*x_value*(predicted_y - y_value)
13         dg_db = 2*(predicted_y - y_value)
14
15         # update
16         a = a - lr*dg_da
17         b = b - lr*dg_db
```

# Summary



$$\frac{dg}{da} = \frac{dg}{df} \frac{df}{da} = 2x(f - y)$$
$$\frac{dg}{db} = \frac{dg}{df} \frac{df}{db} = 2(f - y)$$



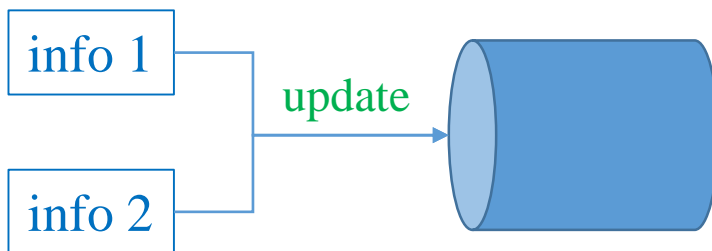


```
1 import random
2
3 # predict function
4 def predicted_func(x, a, b):
5     return a*x + b
6
7 # parameters
8 num_steps = 100
9 lr = 0.01
10
11 # given data
12 x1 = -2
13 y1 = -1
14
15 x2 = 5
16 y2 = 13
17
18 # 1. set a, b randomly
19 a = random.random()*10.0 - 5.0
20 b = random.random()*10.0 - 5.0
```

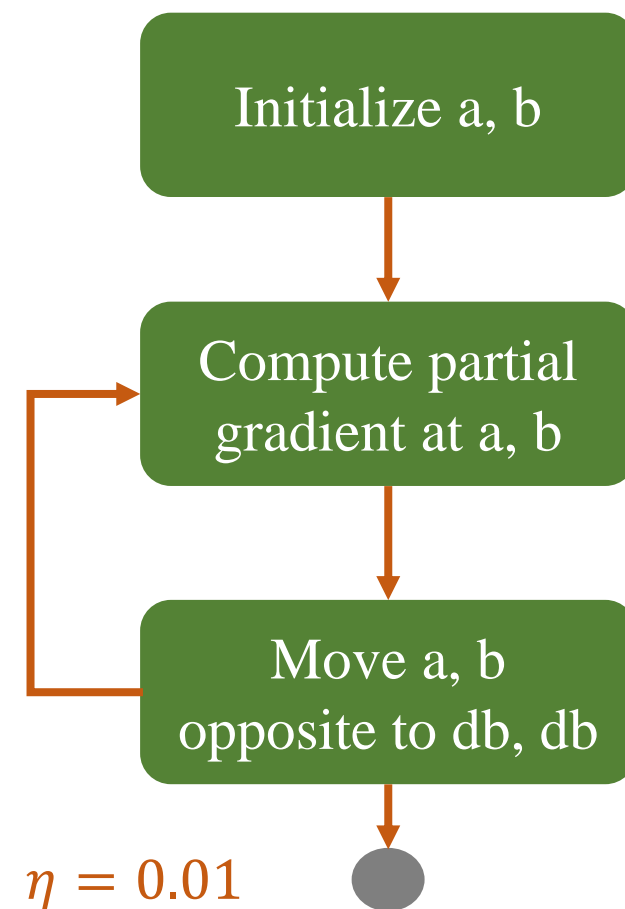
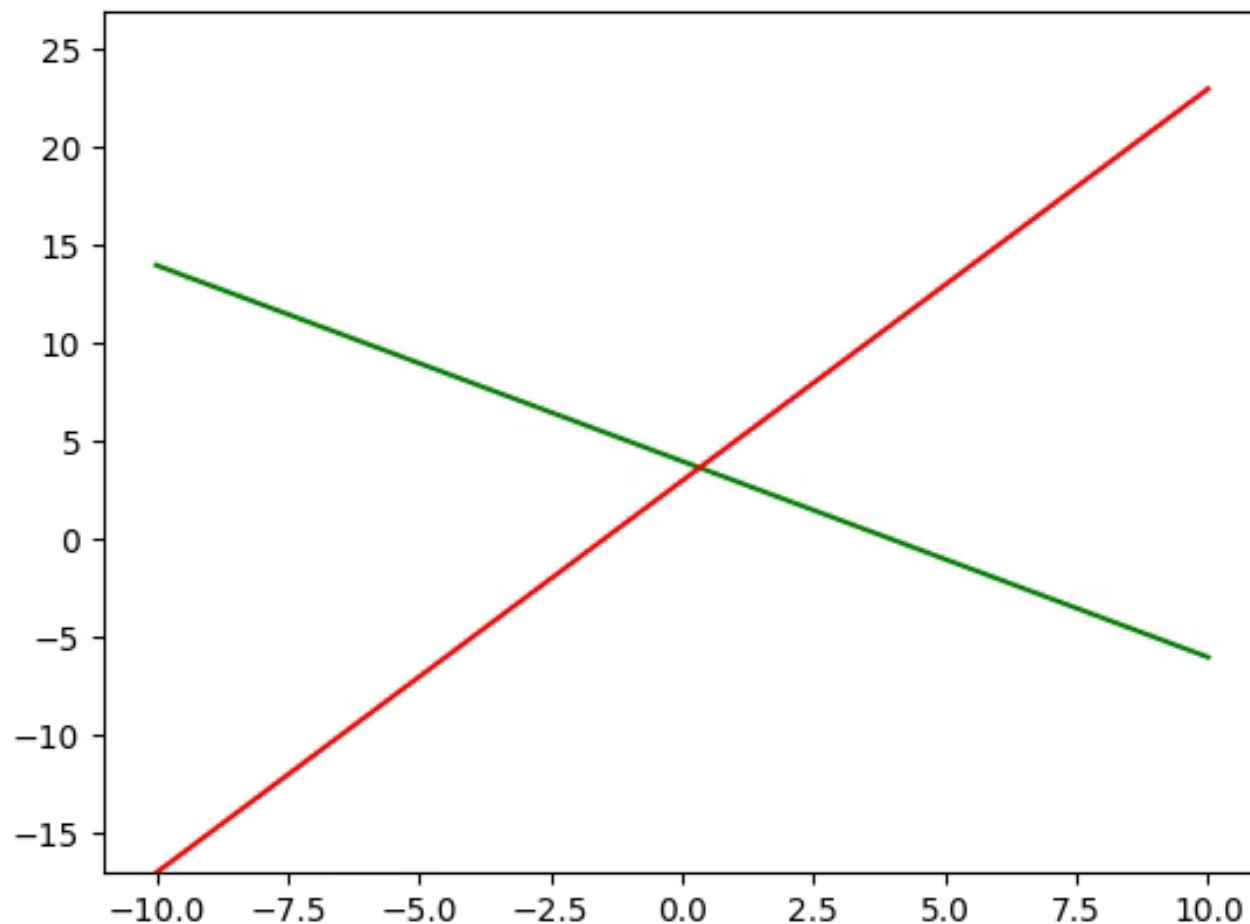
```
1 for i in range(num_steps):
2     # 2. compute predicted_y1 and predicted_y2
3     predicted_y1 = predicted_func(x1, a, b)
4     predicted_y2 = predicted_func(x2, a, b)
5
6     # 3. compute g
7     g_value_1 = (predicted_y1 - y1)**2
8     g_value_2 = (predicted_y2 - y2)**2
9
10    # logging
11    # ...
12
13    # 4. compute partial gradients for a and b
14    dg_da = 2*x1*(predicted_y1 - y1) + 2*x2*(predicted_y2 - y2)
15    dg_db = 2*(predicted_y1 - y1) + 2*(predicted_y2 - y2)
16
17    # 5. update
18    a = a - lr*dg_da
19    b = b - lr*dg_db
```

# Summary

2



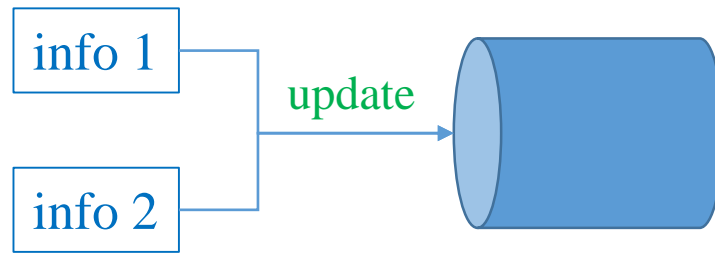
$$\frac{dg}{da} = \frac{dg}{df} \frac{df}{da} = 2x(f - y)$$
$$\frac{dg}{db} = \frac{dg}{df} \frac{df}{db} = 2(f - y)$$



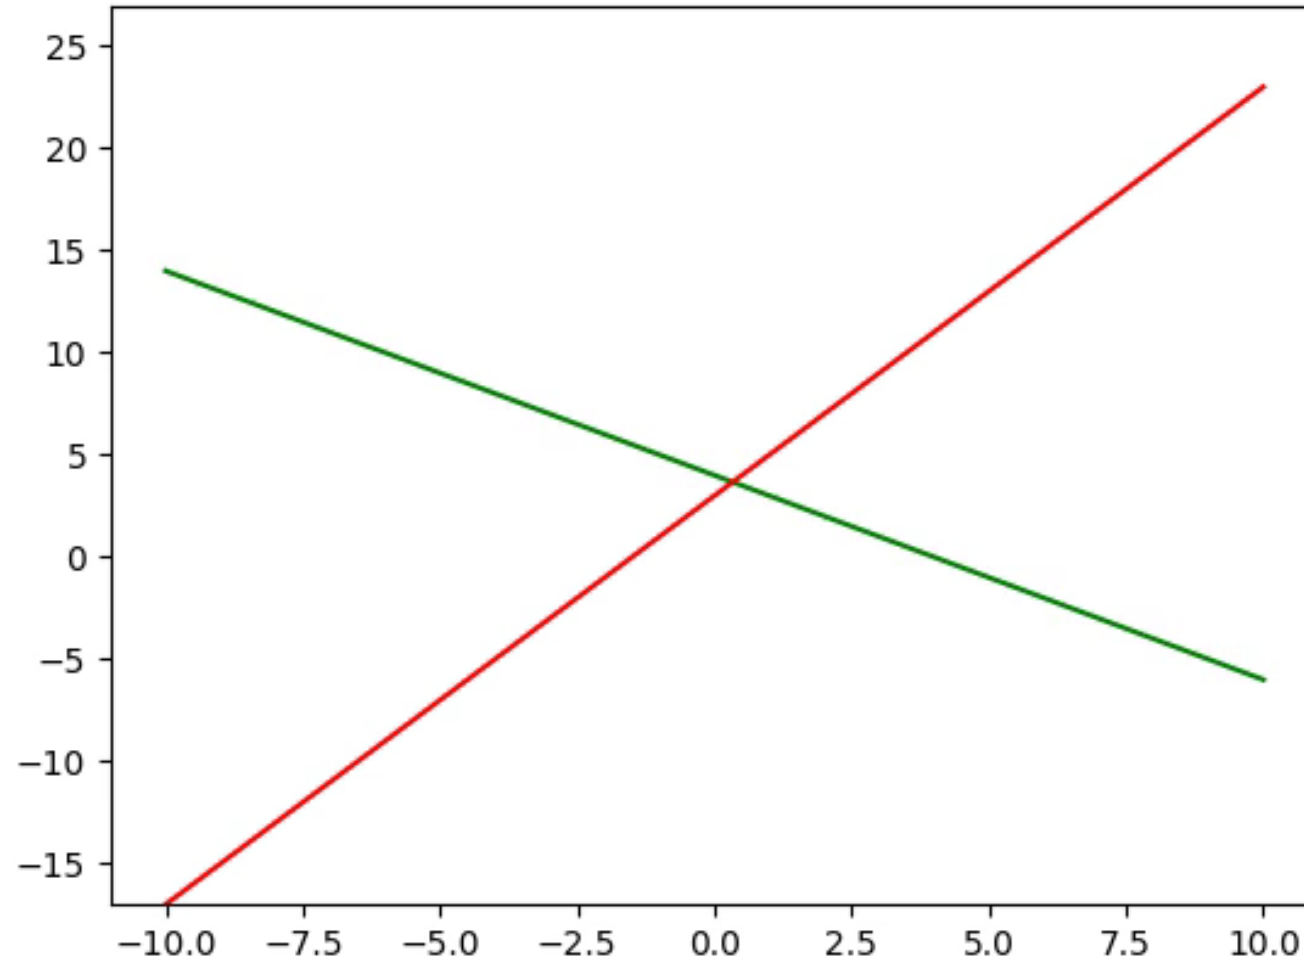


# Summary

2



$$\frac{dg}{da} = \frac{dg}{df} \frac{df}{da} = 2x(f - y)$$
$$\frac{dg}{db} = \frac{dg}{df} \frac{df}{db} = 2(f - y)$$



Initialize a, b

Compute partial  
gradient at a, b

Move a, b  
opposite to db, db

$\eta = 0.001$



# List-based Implementation

## ❖ Vectorization

$$f(x) = ax + b$$

$$f(x) = ax + b$$

$$= ax + b1$$

$$= \begin{bmatrix} x \\ 1 \end{bmatrix} \cdot \begin{bmatrix} a \\ b \end{bmatrix}$$

```
1 # implement 1: naive approach
2 def predict_w1(x, a, b):
3     return a*x + b
4
5 # test
6 print(predict_w1(2, 1, 3))
```

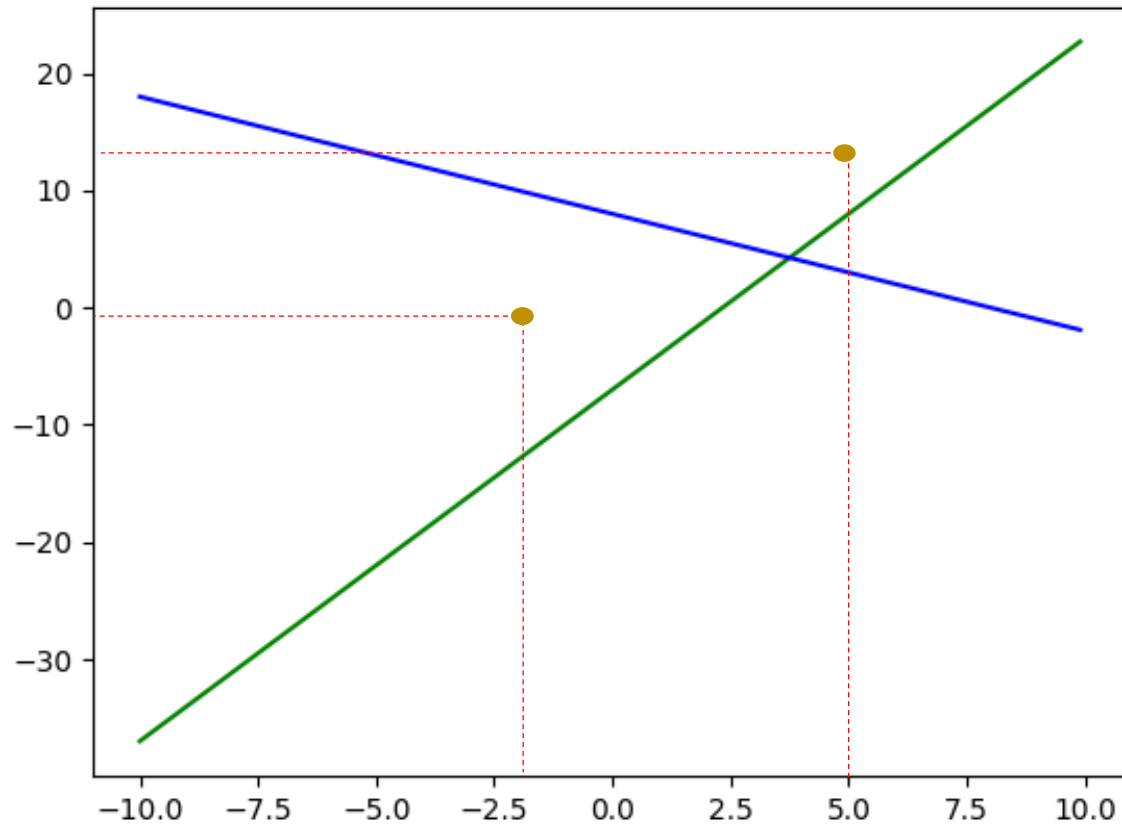
5

```
1 # implement 2: using list
2 def predict_w2(data, weights):
3     return sum([d*w for d, w in zip(data, weights)])
4
5 # test
6 print(predict_w2([2, 1], [1, 3]))
```

5

# Discussion

Remove the red line

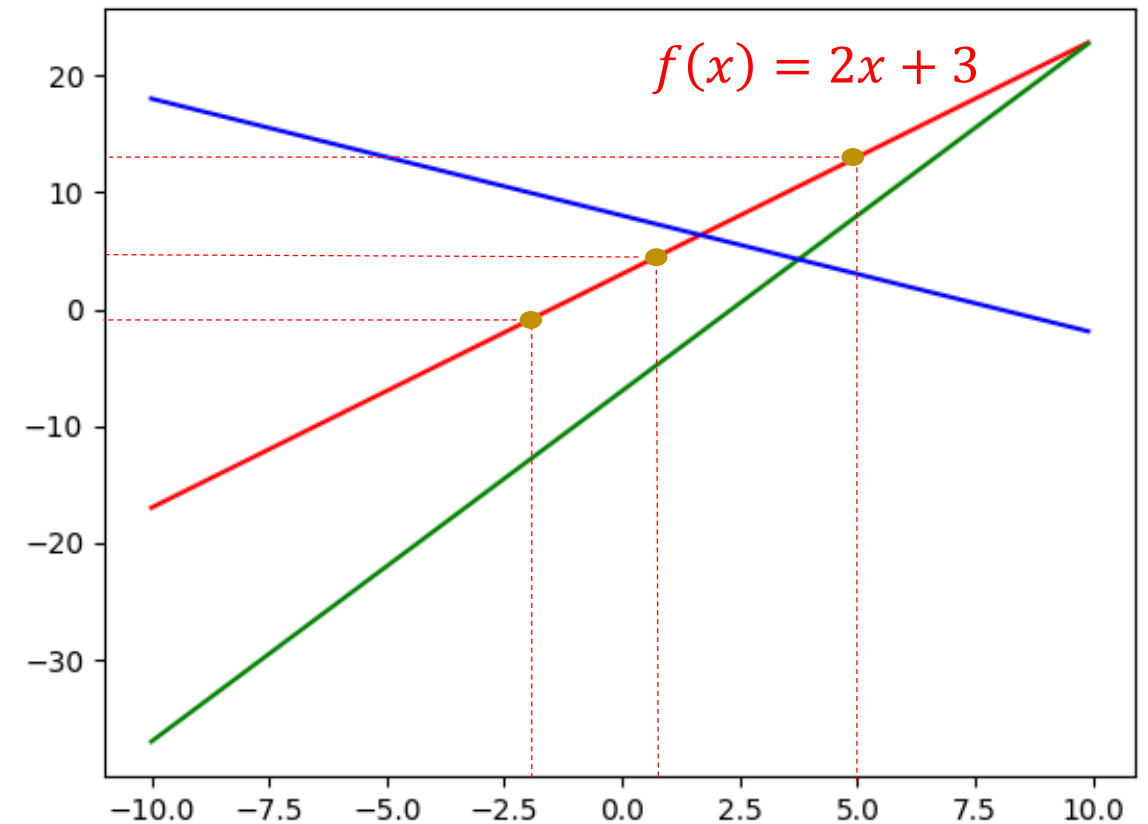


Have one  
more sample

$$(x_1 = -2, y_1 = -1)$$

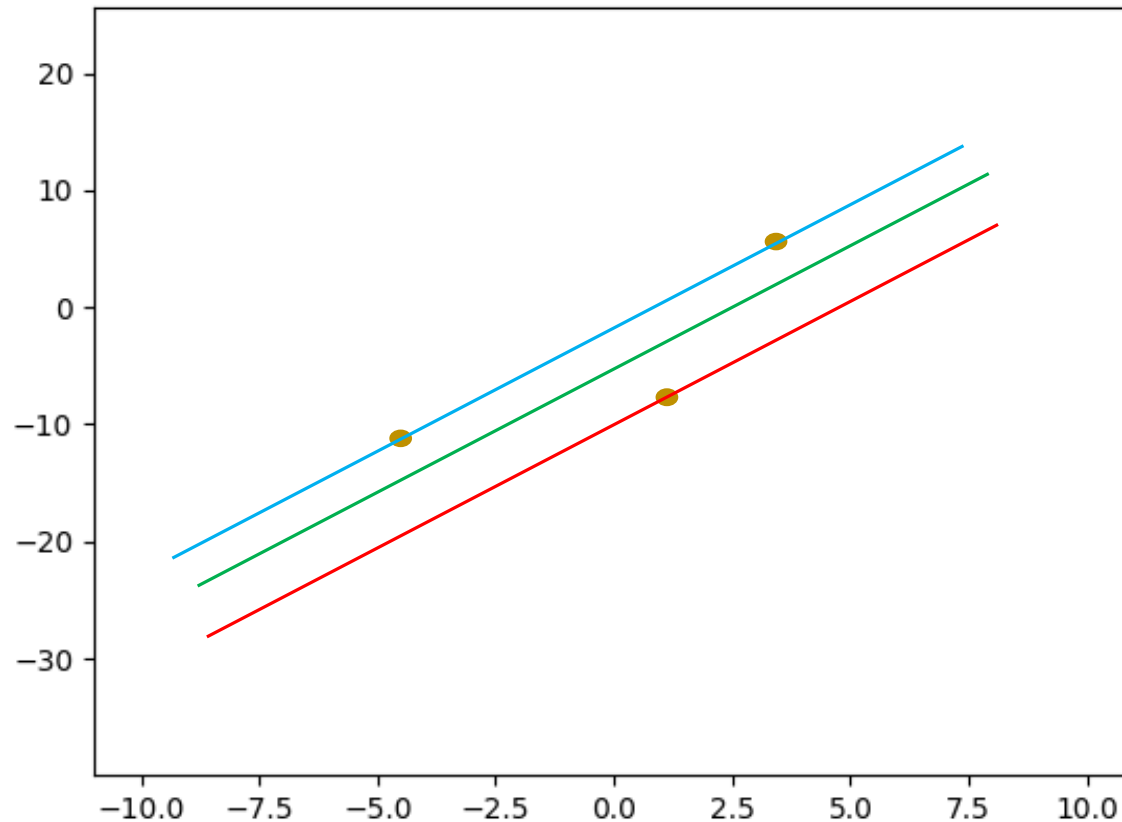
$$(x_2 = 5, y_2 = 13)$$

$$(x_2 = 1, y_2 = 3)$$



# Discussion

❖ What about the given samples?



Line 1: go through two points



Line 2: smallest summation of distances



Line 3: go through one point

